

CCSYA

Instructions: Language of the Computer

Part I

Departamento de Engenharia Informática
Instituto Superior de Engenharia do Porto

Luís Nogueira (lmn@isep.ipp.pt)

Levels of Program Code

■ High-level language

- Level of abstraction closer to problem domain
- Provides for productivity and portability

■ Assembly language

- Textual representation of instructions

■ Hardware representation

- Binary digits (bits)
- Encoded instructions and data

High-level
language
program
(in C)

```
swap(int v[], int k)
{int temp;
  temp = v[k];
  v[k] = v[k+1];
  v[k+1] = temp;
}
```

Compiler

Assembly
language
program
(for MIPS)

```
swap:
    muli $2, $5, 4
    add  $2, $4, $2
    lw   $15, 0($2)
    lw   $16, 4($2)
    sw   $16, 0($2)
    sw   $15, 4($2)
    jr   $31
```

Assembler

Binary machine
language
program
(for MIPS)

```
000000001010000100000000000011000
000000000000110000001100000100001
100011000110001000000000000000000
100011001111001000000000000000100
101011001111001000000000000000000
101011000110001000000000000000100
00000011111000000000000000001000
```

Instruction Set

- To command a computer's hardware, you must speak its language
- The words of a computer's language are called *instructions*, and its vocabulary is called an **instruction set**
- Computer languages are quite similar, more like regional dialects than like independent languages
 - Hence, once you learn one, it is easy to pick up others

Complex Instruction Set (CISC)

- **Primary goal is to complete a task in few lines of assembly instruction as possible**
 - Each instruction can execute several low-level operations, such as a load from memory, an arithmetic operation, and a memory store, all in a single instruction
- **Complexity is on the hardware**
 - Variable length instructions
 - Multiple addressing modes
 - Instructions require multiple clock cycles to execute
 - ...
- **VAX, Motorola 68000, Intel x86, ...**

The Intel x86 ISA

- **Evolution with backward compatibility**
 - 8080 (1974): 8-bit microprocessor
 - Accumulator, plus 3 index-register pairs
 - 8086 (1978): 16-bit extension to 8080
 - Complex instruction set (CISC)
 - 8087 (1980): floating-point coprocessor
 - Adds FP instructions and register stack
 - 80286 (1982): 24-bit addresses, MMU
 - Segmented memory mapping and protection
 - 80386 (1985): 32-bit extension (now IA-32)
 - Additional addressing modes and operations
 - Paged memory mapping as well as segments

The Intel x86 ISA

■ Further evolution...

- i486 (1989): pipelined, on-chip caches and FPU
 - Compatible competitors: AMD, Cyrix, ...
- Pentium (1993): superscalar, 64-bit datapath
 - Later versions added MMX (Multi-Media eXtension) instructions
 - The infamous FDIV bug
- Pentium Pro (1995), Pentium II (1997)
 - New microarchitecture (see Colwell, *The Pentium Chronicles*)
- Pentium III (1999)
 - Added SSE (Streaming SIMD Extensions) and associated registers
- Pentium 4 (2001)
 - New microarchitecture
 - Added SSE2 instructions

The Intel x86 ISA

■ And further...


- AMD64 (2003): extended architecture to 64 bits
- EM64T – Extended Memory 64 Technology (2004)
 - AMD64 adopted by Intel (with refinements)
 - Added SSE3 instructions
- Intel Core (2006)
 - Added SSE4 instructions, virtual machine support
- AMD64 (announced 2007): SSE5 instructions
 - Intel declined to follow, instead...
- Advanced Vector Extension (announced 2008)
 - Longer SSE registers, more instructions

■ If Intel didn't extend with compatibility, its competitors would!

- Technical elegance ≠ market success

Fallacies

- **Powerful instruction \Rightarrow higher performance**
 - Fewer instructions required
 - But complex instructions are hard to implement
 - May slow down all instructions, including simple ones
 - Compilers are good at making fast code from simple instructions
- **Backward compatibility \Rightarrow instruction set doesn't change**
 - But they do accrete more instructions



x86 instruction set

Implementing x86 ISA

- **Complex instruction set makes implementation difficult**
 - Hardware translates instructions to simpler microoperations
 - Simple instructions: 1–1
 - Complex instructions: 1–many
 - Micro engine similar to RISC
 - Market share makes this economically viable
- **Comparable performance to RISC**
 - Compilers avoid complex instructions

x86 ISA

- **What the x86 lacks in style, it made up for in market size, making it beautiful from the right perspective**
 - Its saving grace is that the most frequently used x86 architectural components are not too difficult to implement
- **In the PostPC era, however, despite considerable architectural and manufacturing expertise, x86 has not yet been competitive in the personal mobile device**

Reduced Instruction Set (RISC)

- **A small set of simple and general instructions that can be executed within one clock cycle**
- **Complexity is on the software**
 - Primary goal is to speedup individual instruction execution at the cost of a higher number of instructions in a program
 - A larger number of registers
- **Load/store architecture**
 - Memory is accessed through specific instructions rather than as a part of most instructions
- **ARM, MIPS, IBM PowerPC, Sun Sparc, ...**

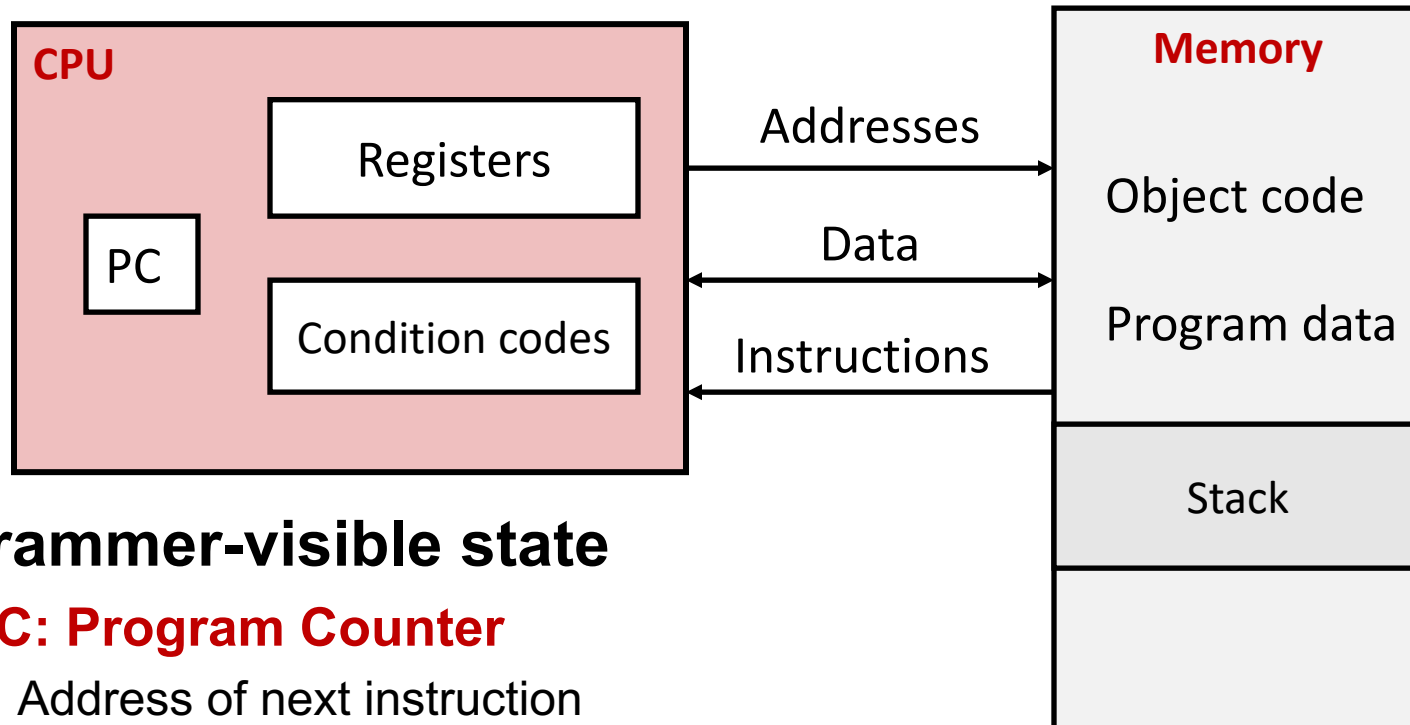
The MIPS Instruction Set

- **Used as the example throughout the course**
- **Stanford MIPS commercialized by MIPS Technologies (www.mips.com)**
- **Similar ISAs have a large share of embedded core market**
 - Applications in consumer electronics, network/storage equipment, cameras, printers, ...
 - ARMv8 is very similar to MIPS

MIPS Simplified Implementation

- **We will consider a simplified implementation of the MIPS instruction set architecture**
 - As a way to illustrate how it determines many aspects of the implementation, and how the choice of different implementation strategies affects the clock rate and CPI for the computer
- **A subset of the core MIPS instruction set:**
 - Arithmetic-logical instructions: **add**, **sub**, **and**, **or**, **slt**, ...
 - Memory reference instructions: **lw**(load word), **sw**(store word)
 - Control transfer instructions: **beq**(branch equal), **bne** (branch not equal), **j**(jump unconditionally), ...

Assembly programmer's view



Programmer-visible state

- **PC: Program Counter**
 - Address of next instruction
- **Registers**
 - Heavily used program data
- **Condition codes**
 - Store status information about most recent arithmetic operation
 - Used for conditional branching

Memory

- Byte addressable array
- Code and user data
- Stack to support procedures

Arithmetic Operations

- **Example: add**

- Two sources and one destination

add a, b, c # a gets b + c

- **All arithmetic operations have this form**

Design Principle 1: Simplicity favors regularity

Regularity makes implementation simpler

Simplicity enables higher performance at lower cost

Register Operands

- **Arithmetic instructions use register operands**
- **MIPS has a 32×32 -bit register file**
 - Use for frequently accessed data
 - Numbered 0 to 31
 - 32-bit data called a “word”
- **Assembler names**
 - `$t0`, `$t1`, ..., `$t9` for temporary values
 - `$s0`, `$s1`, ..., `$s7` for saved variables

Design Principle 2: Smaller is faster

The designer must balance the craving of programs for more registers with the designer’s desire to keep the clock cycle fast

Register Operand Example

- **C code:**

```
f = (g + h) - (i + j);
```

- **Compiled MIPS code:**

- f, ..., j in \$s0, ..., \$s4

```
add $t0, $s1, $s2  
add $t1, $s3, $s4  
sub $s0, $t0, $t1
```

Immediate Operands

- **Constant data specified in an instruction**

```
addi $s3, $s3, 4
```

- **No subtract immediate instruction**

- Just use a negative constant

```
addi $s2, $s1, -1
```

Design Principle 3: Make the common case fast

Small constants are common

Immediate operand avoids a load instruction

Operations are much faster and use less energy than if constants were loaded from memory

The Constant Zero

- **MIPS register 0 (\$zero) is the constant 0**
 - Cannot be overwritten
 - Its main role is to simplify the instruction set by offering useful variations
- **Useful for common operations**
 - E.g., move between registers

```
add $t2, $s1, $zero
```

Bitwise Operations

- **Instructions for bitwise manipulation**
 - Useful for extracting and inserting groups of bits in a word

Operation	C	MIPS
Shift left	<<	sll
Shift right	>>	srl
Bitwise AND	&	and, andi
Bitwise OR		or, ori
Bitwise NOT	~	nor

Shift Operations

■ Shift left logical

- Shift left and fill with 0 bits
- `sll` by i bits multiplies by 2^i

`sll $t2, $s0, 4` # $\$t2 = \$s0 \ll 4$

■ Shift right logical

- Shift right and fill with 0 bits
- `srl` by i bits divides by 2^i (unsigned only)

`srl $t2, $s0, 4` # $\$t2 = \$s0 \gg 4$

AND Operations

- **Useful to mask bits in a word**
 - Select some bits, clear others to 0

and \$t0, \$t1, \$t2

\$t2	0000 0000 0000 0000 0000 1101 1100 0000
\$t1	0000 0000 0000 0000 0011 1100 0000 0000
\$t0	0000 0000 0000 0000 0000 1100 0000 0000

OR Operations

- **Useful to include bits in a word**
 - Set some bits to 1, leave others unchanged

```
or $t0, $t1, $t2
```

\$t2	0000 0000 0000 0000 0000 1101 1100 0000
------	---

\$t1	0000 0000 0000 0000 0011 1100 0000 0000
------	---

\$t0	0000 0000 0000 0000 0011 1101 1100 0000
------	---

NOT Operations

- **Useful to invert bits in a word**
 - Change 0 to 1, and 1 to 0
- **MIPS has NOR 3-operand instruction**
 - $a \text{ NOR } b == \text{NOT} (a \text{ OR } b)$

```
nor $t0, $t1, $zero
```

Register 0: always
read as zero

\$t1 0000 0000 0000 0000 0011 1100 0000 0000

\$t0 1111 1111 1111 1111 1100 0011 1111 1111

Arithmetic and Bitwise Instructions

<code>add</code>	<code>\$1, \$2, \$3</code>	<code>\$s1 = \$2 + \$3</code>	add
<code>addu</code>	<code>\$1, \$2, \$3</code>	<code>\$s1 = \$2 + \$3</code>	add unsigned
<code>addi</code>	<code>\$1, \$2, imm</code>	<code>\$s1 = \$2 + imm</code>	add immediate
<code>addiu</code>	<code>\$1, \$2, imm</code>	<code>\$s1 = \$2 + imm</code>	add immediate unsigned
<code>sub</code>	<code>\$1, \$2, \$3</code>	<code>\$s1 = \$2 - \$3</code>	subtract
<code>and</code>	<code>\$s1, \$s2, \$s3</code>	<code>\$s1 = \$s2 & \$s3</code>	bitwise and
<code>andi</code>	<code>\$s1, \$s2, imm</code>	<code>\$s1 = \$s2 & \$s3</code>	bitwise and immediate
<code>or</code>	<code>\$s1, \$s2, \$s3</code>	<code>\$s1 = \$s2 \$s3</code>	bitwise or
<code>nor</code>	<code>\$s1, \$s2, \$s3</code>	<code>\$s1 = ~(\$s2 \$s3)</code>	bitwise not
<code>sll</code>	<code>\$s1, \$s2, imm</code>	<code>\$s1 = \$s2 << imm</code>	shift left logical
<code>srl</code>	<code>\$s1, \$s2, imm</code>	<code>\$s1 = \$s2 >> imm</code>	shift right logical

Registers vs. Memory


- **Registers are faster to access than memory**
- **Operating on memory data requires loads and stores**
 - More instructions to be executed
- **Compiler must use registers for variables as much as possible**
 - Only spill to memory for less frequently used variables
 - Register optimization is important!

Memory Operands

- **Main memory used for composite data**
 - Arrays, structures, dynamic data
- **Memory is byte addressed**
 - Just a large, single-dimensional array, with the address acting as the index to that array
 - Words are aligned in memory (address must be a multiple of 4)
- **MIPS is Big Endian**
 - Most-significant byte at least address of a word
 - *c.f.* Little Endian: least-significant byte at least address
 - SPIM simulator's byte order is the same as the byte order of the underlying machine that runs the simulator

Accessing memory

- **Arithmetic operations occur only on registers**
- **Data transfer instructions**
 - Transfer data between memory and registers
- **lw/sw instructions must supply the memory address to access a word in memory**



Addressing Modes

- **Base or displacement addressing**

- The operand is at the memory location whose address is the sum of a register and a constant in the instruction

- **Immediate addressing**

- The operand is a constant in the instruction itself (for jumps and conditional branches)


Memory Operand Example 1

■ C code:

```
g = h + A[8];
```

■ Compiled MIPS code:

- g in \$s1, h in \$s2, base address of A in \$s3
- Index 8 requires offset of 32
 - 4 bytes per word



Memory Operand Example 2

- **C code:**

```
A[12] = h + A[8];
```

- **Compiled MIPS code:**

- h in \$s2, base address of A in \$s3
- Index 8 requires offset of 32

```
lw    $t0, 32($s3)    # load word
add   $t0, $s2, $t0
sw    $t0, 48($s3)    # store word
```

Byte/Halfword Operations

`lb rt, offset(rs)`

- Sign extend to 32 bits in rt

`lh rt, offset(rs)`

`lbu rt, offset(rs)`

- Zero extend to 32 bits in rt

`lhu rt, offset(rs)`

`sb rt, offset(rs)`

- Store just rightmost byte/halfword

`sh rt, offset(rs)`

Data Transfer Instructions

lw \$s1, imm(\$s2)	\$s1 = MEM[\$2 + imm]	load word from memory
lh \$s1, imm(\$s2)	\$s1 = MEM[\$2 + imm]	load half word (2 bytes)
lhu \$s1, imm(\$s2)	\$s1 = MEM[\$2 + imm]	load half word unsigned
lb \$s1, imm(\$s2)	\$s1 = MEM[\$2 + imm]	load byte
lbu \$s1, imm(\$s2)	\$s1 = MEM[\$2 + imm]	load byte unsigned
li \$s1, imm	\$s1 = imm	load immediate
la \$s1, L	\$s1 = L	load address
sw \$s1, imm(\$s2)	MEM[\$2 + imm] = \$1	store word to memory
sh \$s1, imm(\$s2)	MEM[\$2 + imm] = \$1	store half word (2 bytes)
sb \$s1, imm(\$s2)	MEM[\$2 + imm] = \$1	store byte
move \$s1, \$2	\$s1 = \$2	data move

Conditional Operations

- **Branch to a labeled instruction if a condition is true**
 - Otherwise, continue sequentially
- **`beq rs, rt, L1`**
 - If $(rs == rt)$ branch to instruction labeled L1
- **`bne rs, rt, L1`**
 - If $(rs != rt)$ branch to instruction labeled L1
- **Notice that the assembler relieves the compiler and the assembly language programmer from the tedium of calculating addresses for branches**
 - Just as it does for calculating data addresses for loads and stores


Jump Instructions

- There is another kind of branch, often called an **unconditional branch**
 - These instructions say that the processor always follows the branch
- The MIPS name for this type of instruction is **jump**
 - To distinguish between conditional and unconditional branches,
- **j L1**
 - Unconditional jump to instruction labeled L1

Compiling If Statements

■ C code:

```
if (i==j)
    f = g+h;
else
    f = g-h;
```



■ Compiled MIPS code:

- `f, g, ...` in `$s0, $s1, ...`

```
        bne $s3, $s4, Else
        add $s0, $s1, $s2
        j   Exit
Else:    sub $s0, $s1, $s2
Exit:    ...
```

Assembler calculates addresses

Compiling Loop Statements

- **C code:**

```
while (save[i] == k)
    i += 1;
```

- **Compiled MIPS code:**

- i in \$s3, k in \$s5, address of save in \$s6

```
Loop: sll    $t1, $s3, 2
      add    $t1, $t1, $s6
      lw     $t0, 0($t1)
      bne    $t0, $s5, Exit
      addi   $s3, $s3, 1
      j      Loop
Exit: ...
```

More Conditional Operations

- **Set result to 1 if a condition is true**
 - Otherwise, set to 0
- **`slt rd, rs, rt`**
 - if ($rs < rt$) $rd = 1$; else $rd = 0$;
- **`slti rt, rs, constant`**
 - if ($rs < \text{constant}$) $rt = 1$; else $rt = 0$;
- **Usually used in combination with `beq`, `bne`**

```
slt $t0, $s1, $s2    #if ($s1 < $s2)
bne $t0, $zero, L     # branch to L
```

Branch Instruction Design

- Why not b1t, bge, etc in hardware?
- Hardware for $<$, \geq , ... slower than $=$, \neq
 - Combining with branch involves more work per instruction, requiring a slower clock
 - All instructions penalized!
- beq and bne are the common case
- Another example of the great idea of **making the common case fast**

Assembler Pseudo-instructions

- **Most assembler instructions represent machine instructions one-to-one**
- **Pseudo-instructions**
 - The assembler can also treat common variations of machine instructions as if they were instructions in their own right
 - The hardware does not implement the pseudo-instructions, but their appearance in assembly language simplifies programming

```
move $t0, $t1  
li $s1, 20  
blt $t0, $t1, L
```



```
add $t0, $zero, $t1  
addiu $s1, $zero, 20  
slt $at, $t0, $t1  
bne $at, $zero, L
```

\$at (register 1): assembler temporary

Branch Instructions

beq \$1,\$2,L	if(\$s1 == \$2) goto L	branch on equal
bne \$1,\$2,L	if(\$s1 != \$2) goto L	branch on not equal
blt \$1,\$2,L	if(\$s1 < \$2) goto L	branch on less than
bgt \$1,\$2,L	if(\$s1 > \$2) goto L	branch on greater than
ble \$1,\$2,L	if(\$s1 <= \$2) goto L	branch on less than or equal
bge \$1,\$2,L	if(\$s1 >= \$2) goto L	branch on greater than or equal

Jump Instructions

j	imm	goto imm	jump to target address
j	L	goto L	jump to target label
jal	L	\$ra=PC+4; goto L	jump and link (for procedure call)
jr	\$ra	goto \$ra	jump register (for procedure return)

Concluding Remarks

- **Above the machine level is assembly language, a language that humans can read**
 - Each category of MIPS instructions is associated with constructs that appear in high-level programming languages
- **The assembler translates it into the binary numbers that machines can understand...**
- **...and it even “extends” the instruction set by creating symbolic instructions that aren’t in the hardware**
 - Hiding details from the higher level is another example of the great idea of abstraction