

CCSYA

The Processor

Pipelining – Part II

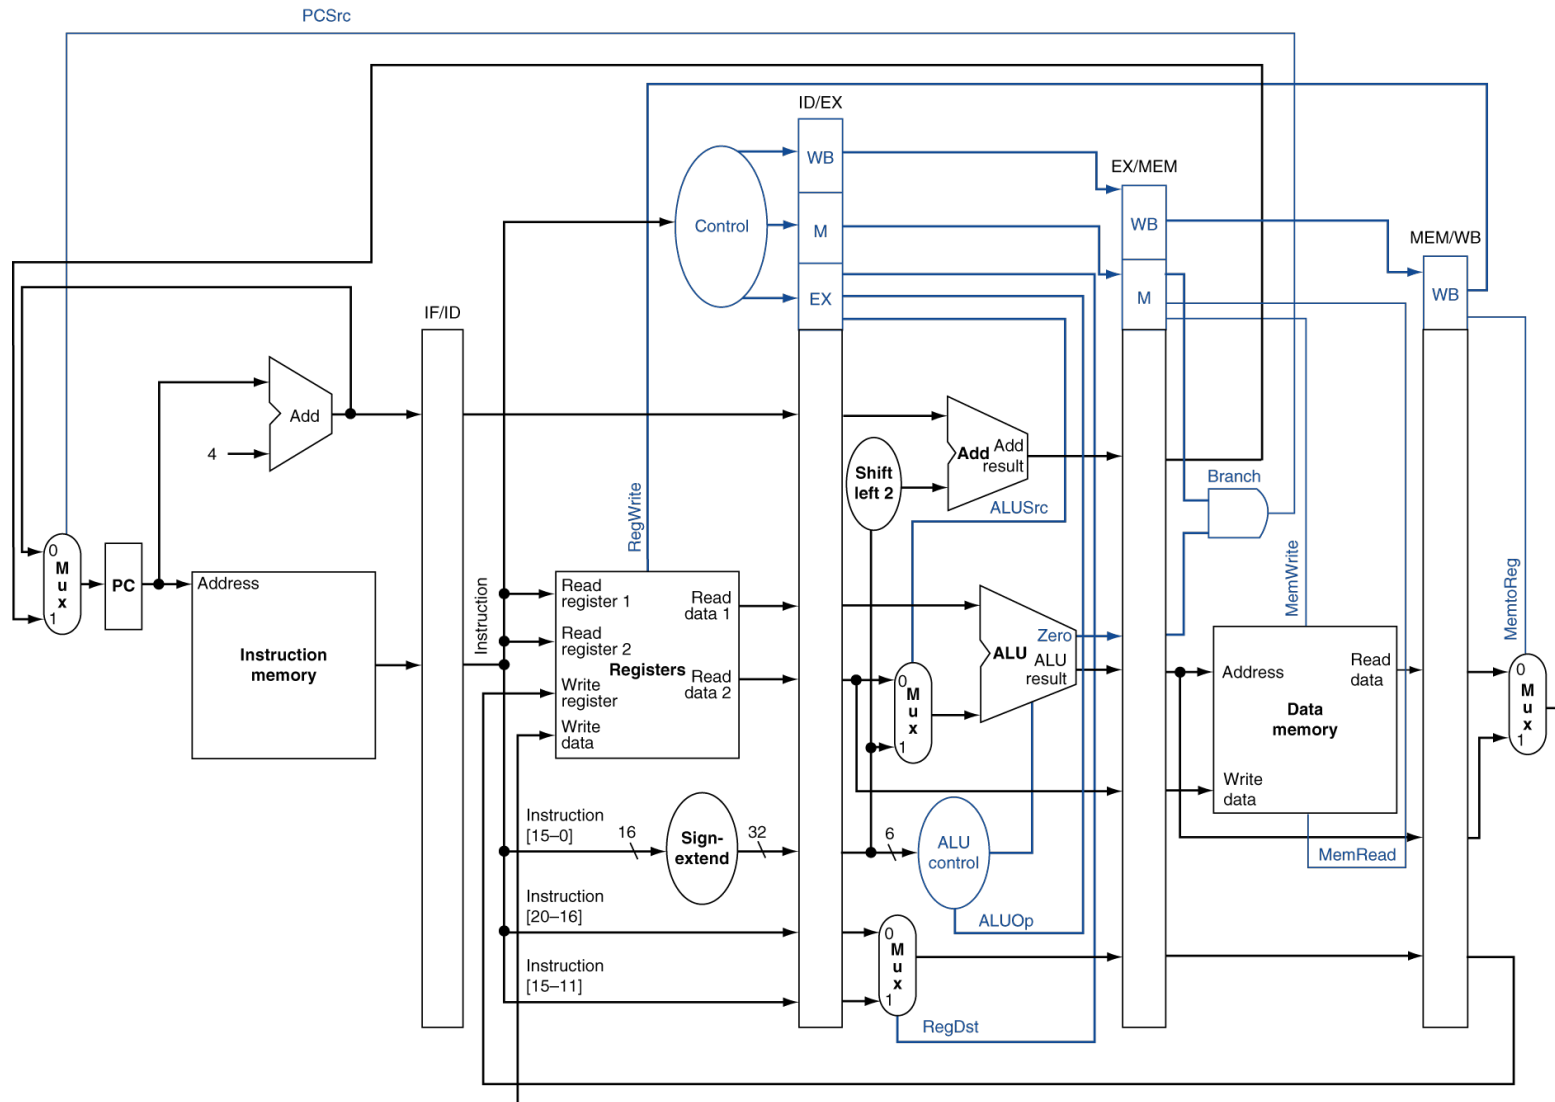
Departamento de Engenharia Informática
Instituto Superior de Engenharia do Porto

Luís Nogueira (lmn@isep.ipp.pt)

Pipeline Summary

- **Pipelining improves performance by increasing instruction throughput**
 - Executes multiple instructions in parallel
 - Each instruction has the same latency
- **Instruction set design affects complexity of pipeline implementation**
 - MIPS ISA designed for pipelining

Pipelined Datapath



Pipeline Hazards

- There are situations when the next instruction cannot execute in the following clock cycle
 - These events are called **hazards**
- **Structural hazard**
 - A required resource is busy
- **Data hazard**
 - Need to wait for previous instruction to complete its data read/write
- **Control hazard**
 - Deciding on control action depends on previous instruction

Dealing with Hazards

- **Without intervention, pipeline must always be stalled**
 - Stalling the pipeline usually lets some instruction(s) proceed, another/others wait for data, resource, etc
- **Stalls impede progress of a pipeline and result in deviation from 1 instruction executing/clock cycle**
 - If no stalls, speedup equal to number of pipeline stages in ideal case
- **CPI pipelined = Ideal CPI + Pipeline stall cycles per instruction**
 - $1 + \text{Pipeline stall cycles per instruction}$

Structural Hazards

- Conflict for use of a resource
- If our MIPS pipeline had a single memory
 - Load/store requires data access
 - Instruction fetch would have to stall for that cycle
 - Would cause a pipeline “bubble”

	← Clock Number →									
Inst. #	1	2	3	4	5	6	7	8	9	10
LOAD	IF	ID	EX	MEM	WB					
Inst. i+1		IF	ID	EX	MEM	WB				
Inst. i+2			IF	ID	EX	MEM	WB			
Inst. i+3				stall	IF	ID	EX	MEM	WB	
Inst. i+4						IF	ID	EX	MEM	WB
Inst. i+5							IF	ID	EX	MEM
Inst. i+6								IF	ID	EX

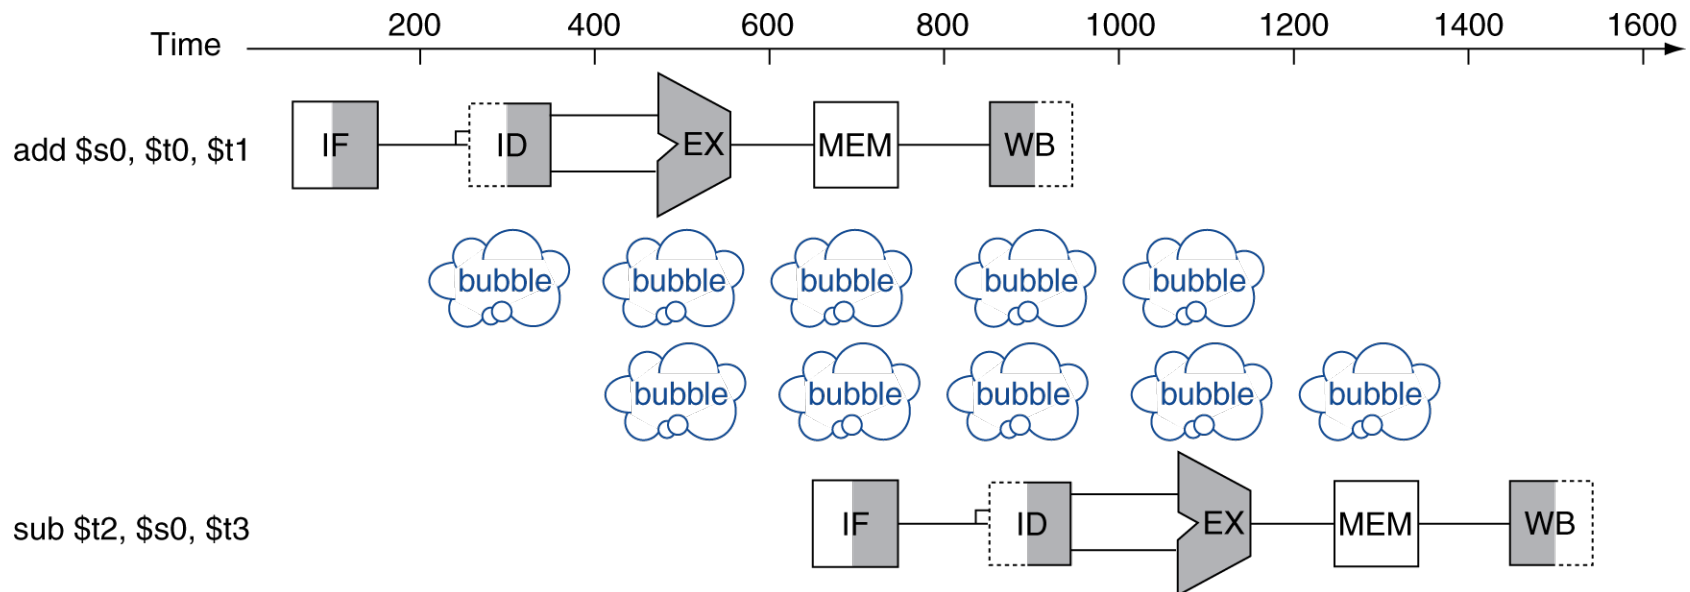
Structural Hazards

- **Solution is to replicate resource**
 - Otherwise, CPI degrades quickly from our ideal '1' for even the simplest of cases
- **MIPS instruction set was designed to be pipelined**
 - Making it easy for designers to avoid structural hazards when designing a pipeline
- **Examples**
 - Separate instruction/data memories
 - Separate instruction/data caches
 - An ALU to perform an arithmetic operation and an adder to increment PC

Data Hazards

- An instruction depends on completion of data access by a previous instruction

- add **\$s0**, \$t0, \$t1
 sub \$t2, **\$s0**, \$t3



Data Hazards

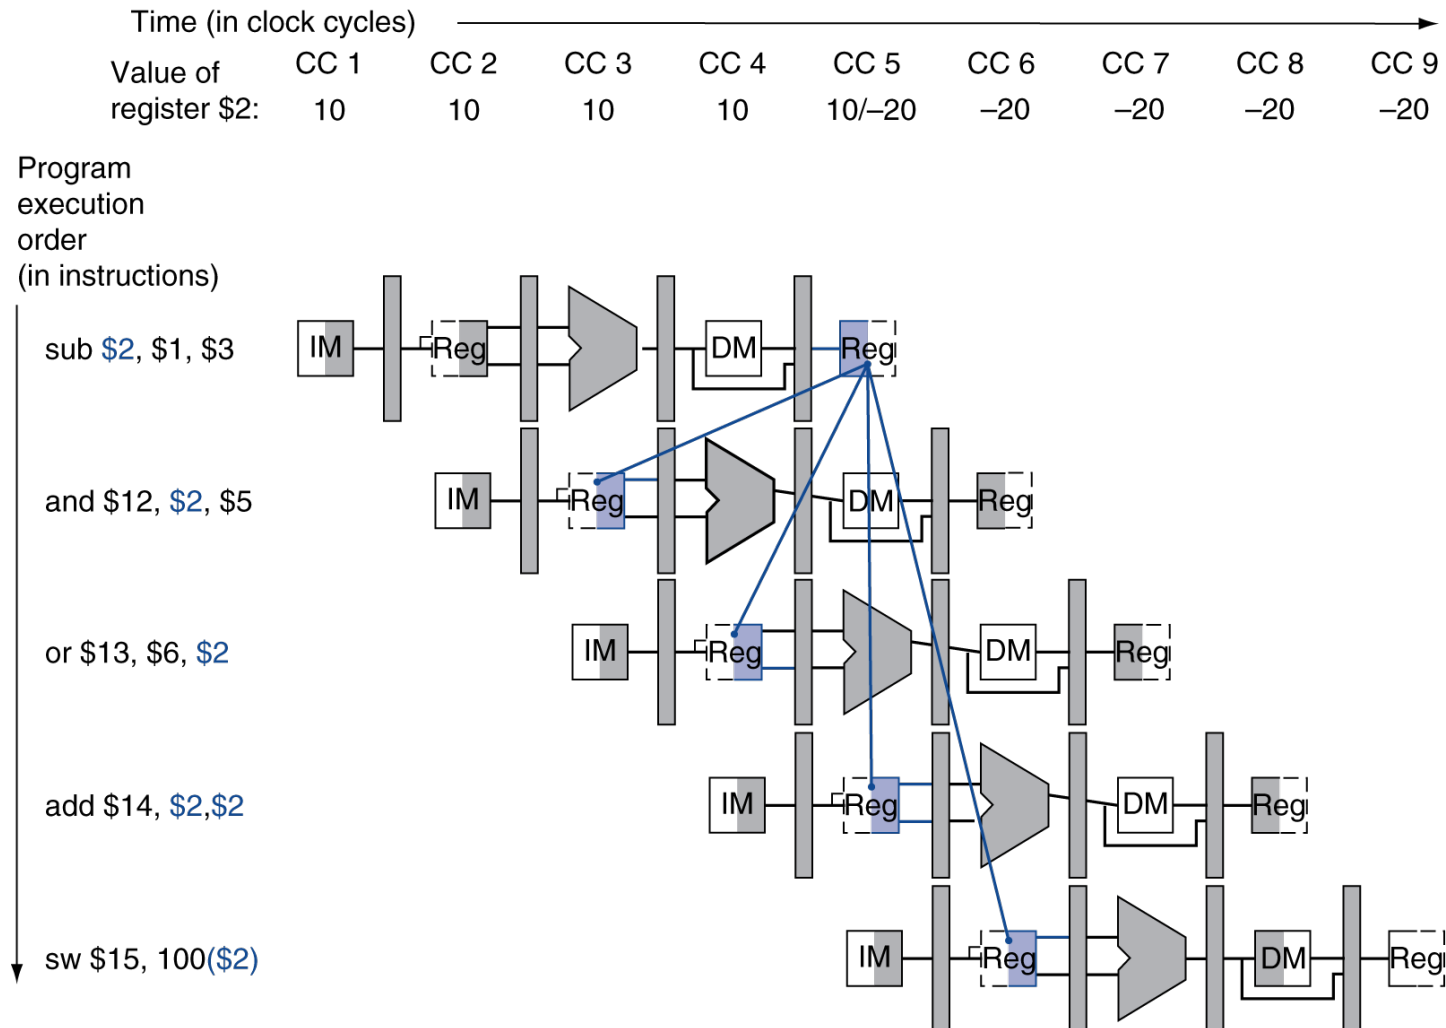
- **Dependence of instructions**

- When an instruction tries to use a register in its EX stage that an earlier instruction intends to write in its WB stage

- **Let's look at a sequence with many dependences**

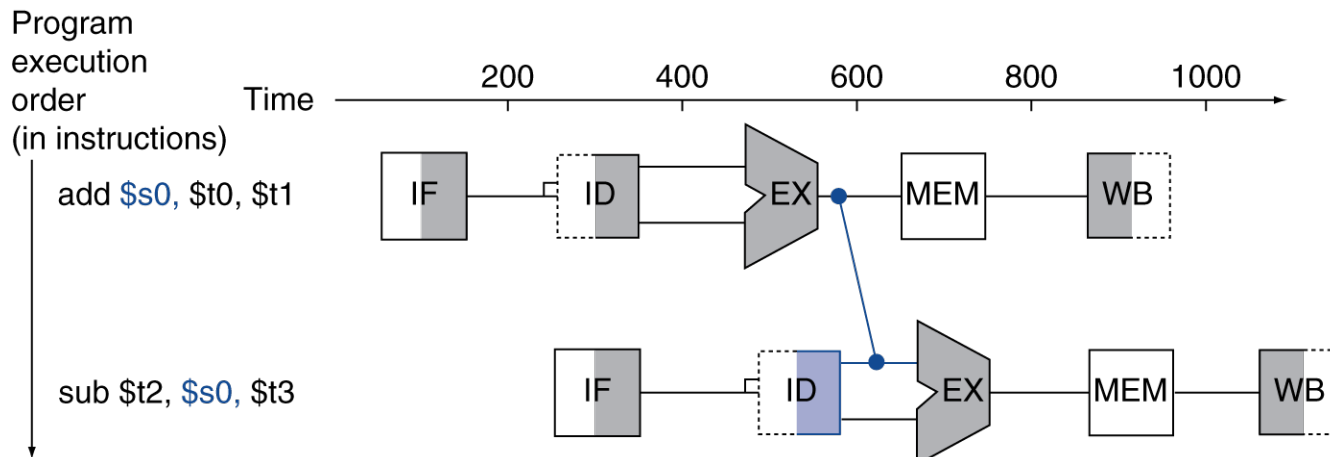
sub	\$2 , \$1, \$3	# Register \$2 written by sub
and	\$12, \$2 , \$5	# 1st operand(\$2) depends on sub
or	\$13, \$6, \$2	# 2nd operand(\$2) depends on sub
add	\$14, \$2, \$2	# 1st(\$2) & 2nd(\$2) depend on sub
sw	\$15, 100(\$2)	# Base (\$2) depends on sub

Data Hazards



Forwarding (aka Bypassing)

- Use result **when it is computed**
 - Don't wait for it to be stored in a register
 - Requires extra connections in the datapath



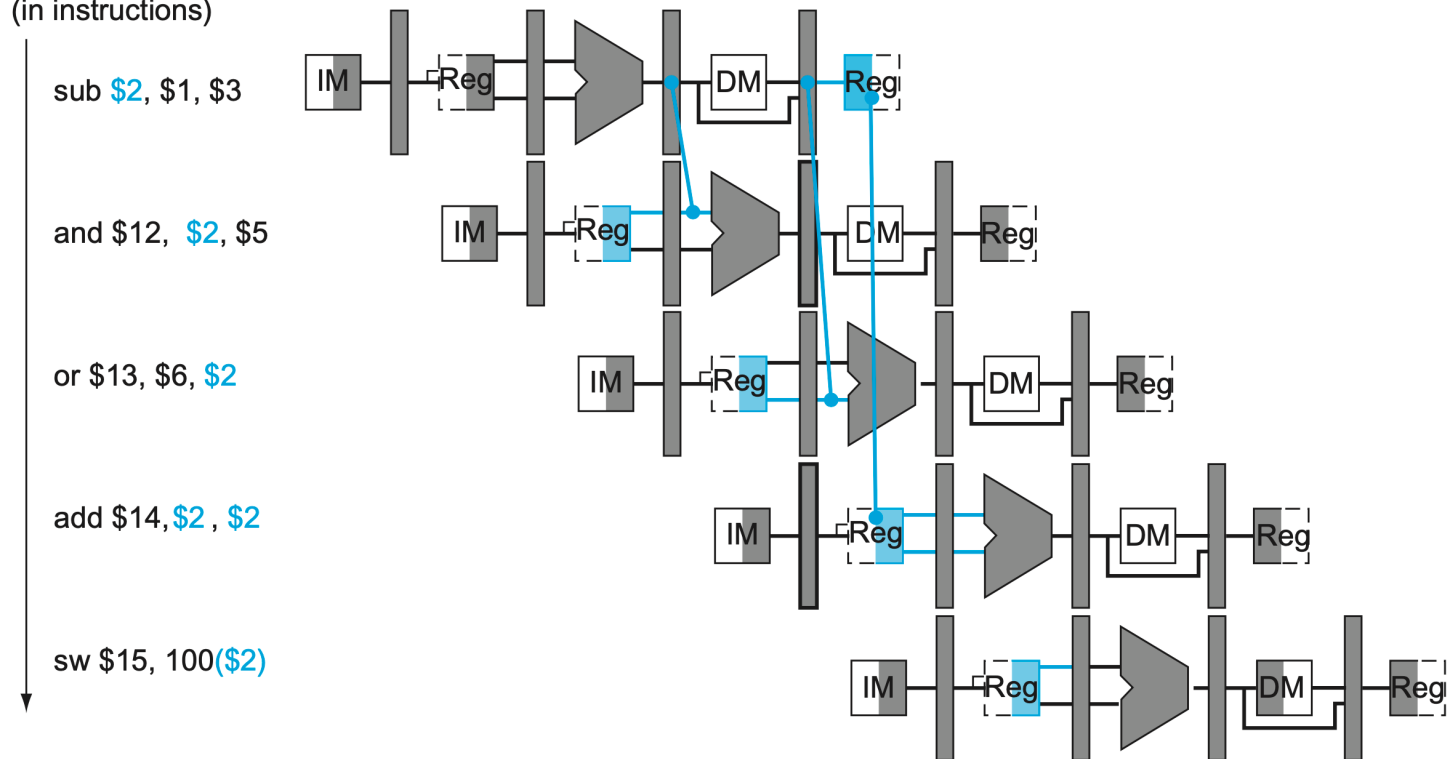
Forwarding

- Dependence is **solved from a pipeline register**, rather than waiting for the WB stage to write the register file
 - Thus, the required data exists in time for later instructions, with the pipeline registers holding the data to be forwarded
- If we can take the inputs to the ALU **from any pipeline register rather than just ID/EX**, then we can forward the proper data
 - By adding multiplexors to the input of the ALU, and with the proper controls, we can run the pipeline at full speed in the presence of these data dependences
 - This forwarding control will be in the EX stage, because the ALU forwarding multiplexors are found in that stage

Forwarding

	Time (in clock cycles) →								
	CC 1	CC 2	CC 3	CC 4	CC 5	CC 6	CC 7	CC 8	CC 9
Value of register \$2:	10	10	10	10	10/-20	-20	-20	-20	-20
Value of EX/MEM:	X	X	X	-20	X	X	X	X	X
Value of MEM/WB:	X	X	X	X	-20	X	X	X	X

Program
execution
order
(in instructions)



Detecting the Need to Forward

- **Pass register numbers along pipeline**
 - e.g., ID/EX.RegisterRs = register number for Rs sitting in ID/EX pipeline register
- **ALU operand register numbers in EX stage are given by ID/EX.RegisterRs, ID/EX.RegisterRt**
- **Data hazards when**
 - 1a. EX/MEM.RegisterRd = ID/EX.RegisterRs
 - 1b. EX/MEM.RegisterRd = ID/EX.RegisterRt
 - 2a. MEM/WB.RegisterRd = ID/EX.RegisterRs
 - 2b. MEM/WB.RegisterRd = ID/EX.RegisterRt

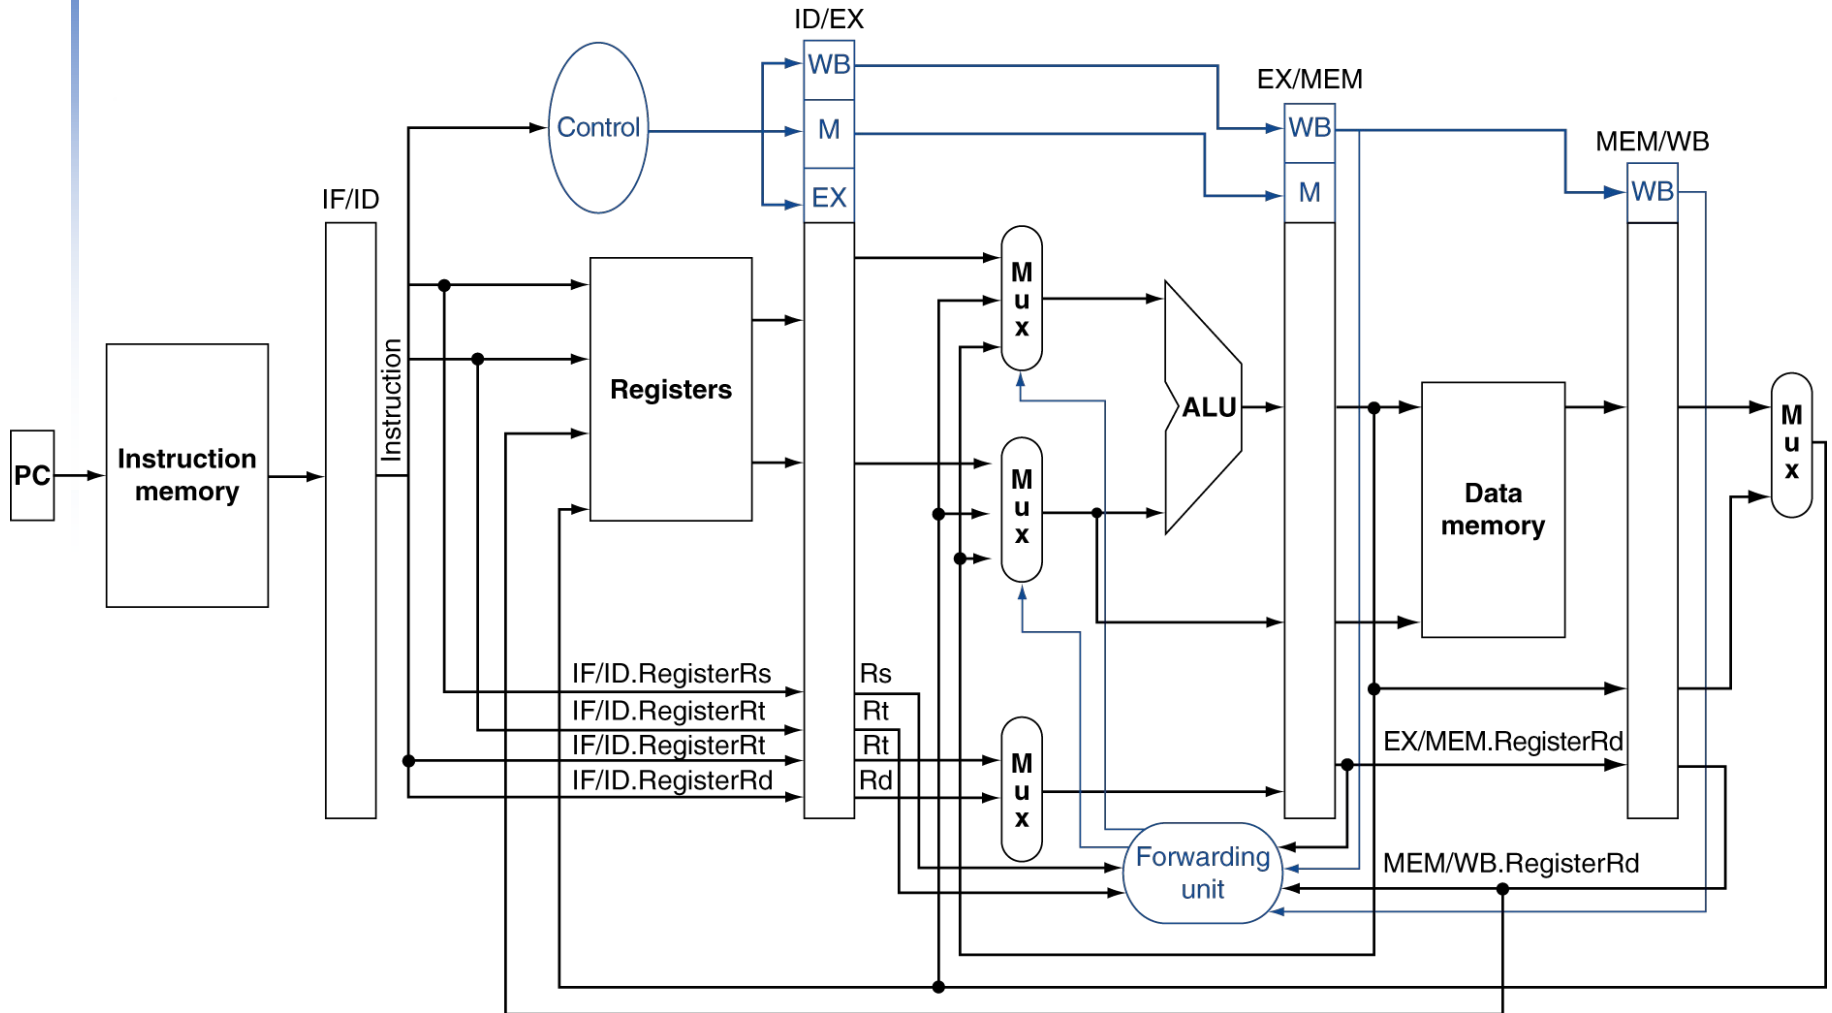
Forward from EX/MEM
pipeline register

Forward from MEM/WB
pipeline register

Detecting the Need to Forward

- **Some instructions do not write registers**
 - No need to forward!
- **One solution is simply to check to see if the RegWrite control signal is active**
 - Examining the WB control field of the pipeline register during the EX and MEM stages determines whether RegWrite is asserted
- **And only if Rd for that instruction is not \$zero**
 - $\text{EX/MEM.RegisterRd} \neq 0$, $\text{MEM/WB.RegisterRd} \neq 0$

Datapath with Forwarding

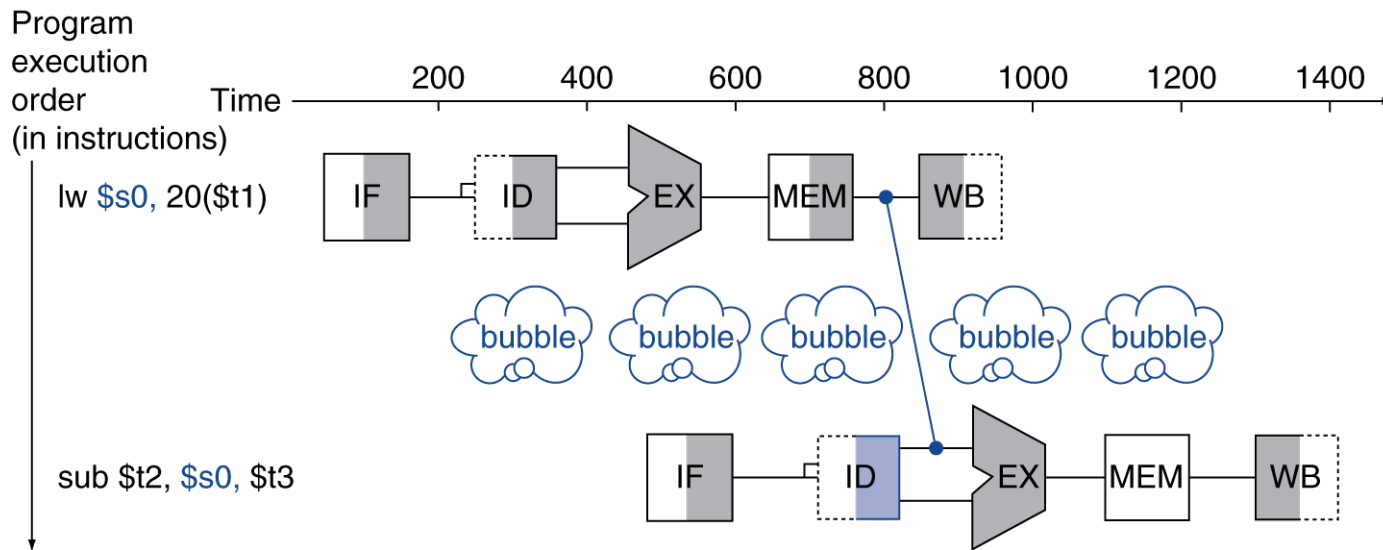


Forwarding Conditions

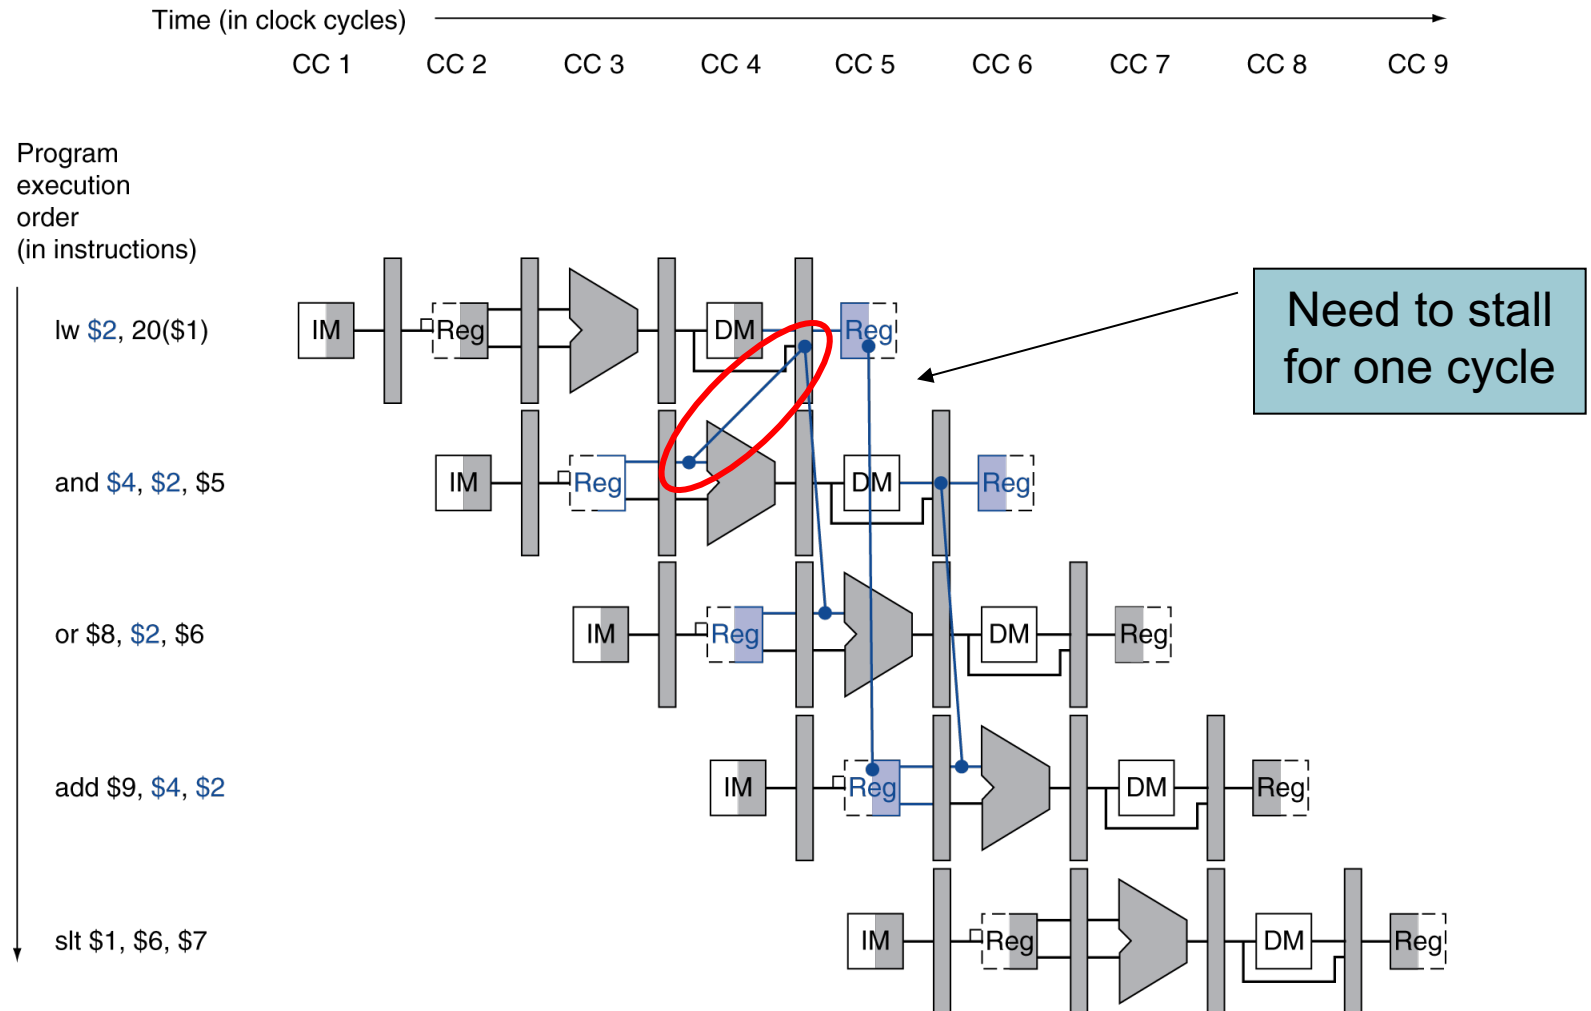
Mux control	Source	Explanation
ForwardA = 00	ID/EX	The first ALU operand comes from the register file
ForwardA = 10	EX/MEM	The first ALU operand is forwarded from the prior ALU result
ForwardA = 01	MEM/WB	The first ALU operand is forwarded from data memory or an earlier ALU result
ForwardB = 00	ID/EX	The second ALU operand comes from the register file
ForwardB = 10	EX/MEM	The second ALU operand is forwarded from the prior ALU result
ForwardB = 01	MEM/WB	The second ALU operand is forwarded from data memory or an earlier ALU result

Load-Use Data Hazard

- **Can't always avoid stalls by forwarding**
 - If value not computed when needed
 - Can't forward backward in time!
- **Load has a latency that forwarding can't solve**



Load-Use Data Hazard



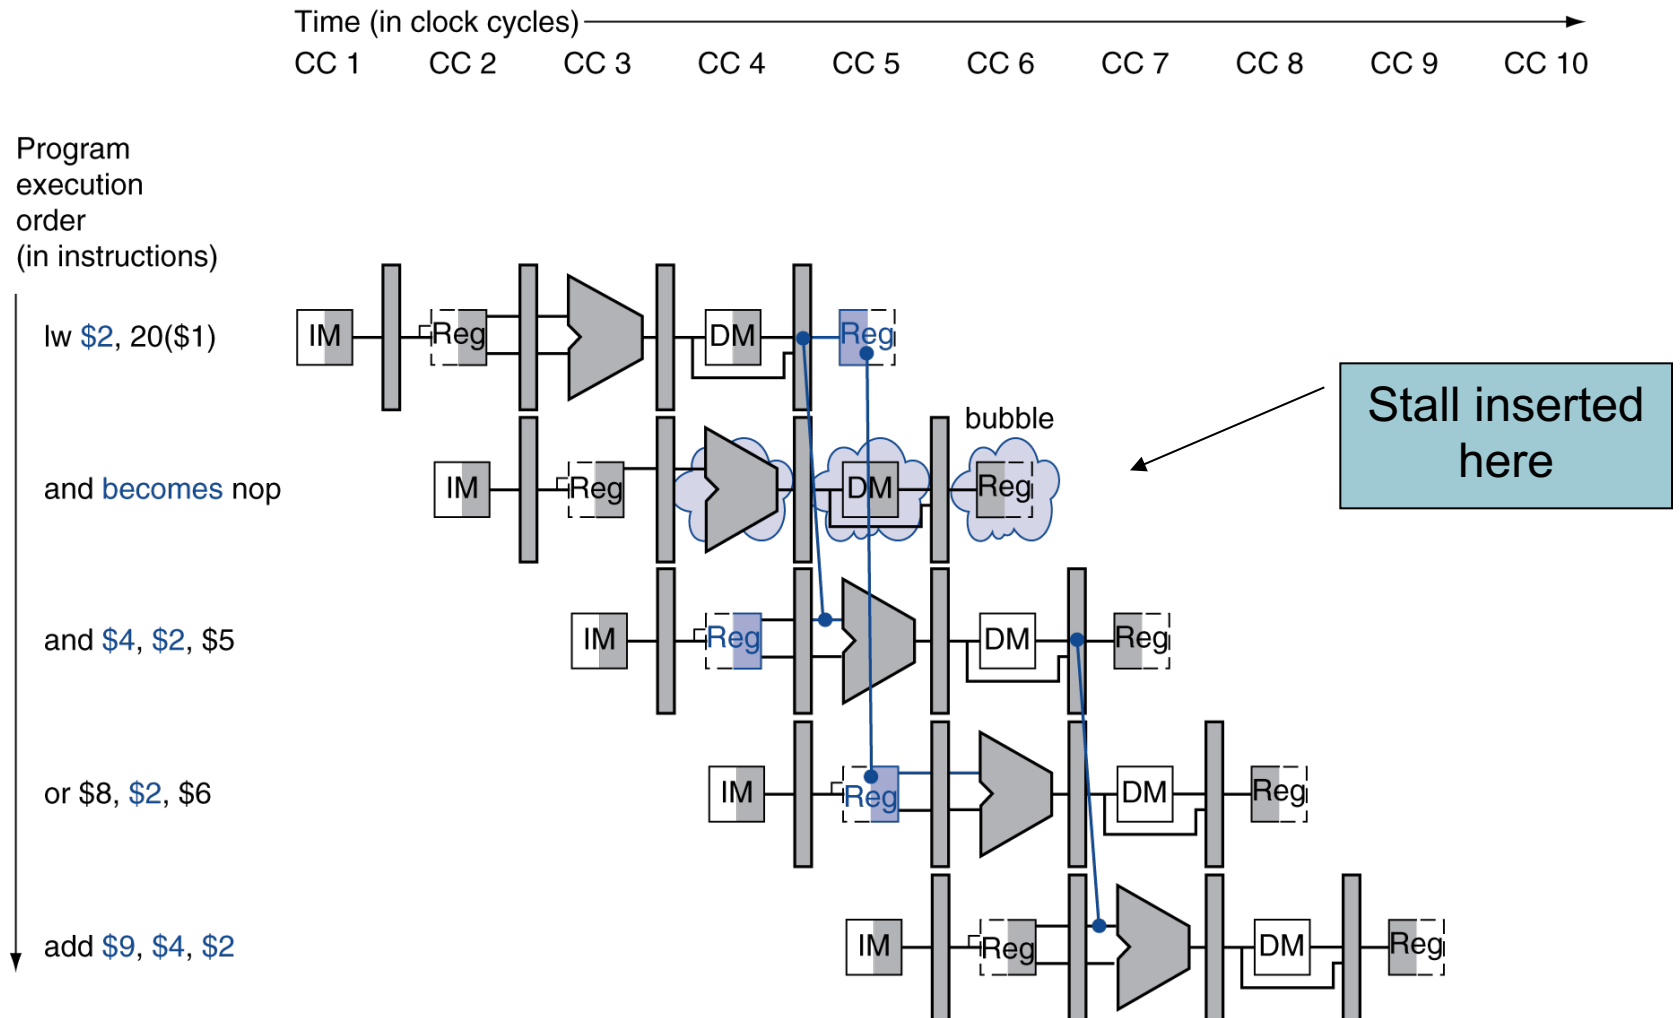
Load-Use Hazard Detection

- We need a **hazard detection unit** operating during the ID stage
 - So that it can insert the stall between the load instruction and its use
- The hazard detection unit must check
 - If a load is in execution and...
 - ...if the destination register of the load in the EX stage matches either source register of the instruction in the ID stage
- Stall the pipeline one clock cycle
 - if (ID/EX.MemRead and ((ID/EX.RegisterRt = IF/ID.RegisterRs) or (ID/EX.RegisterRt = IF/ID.RegisterRt)))

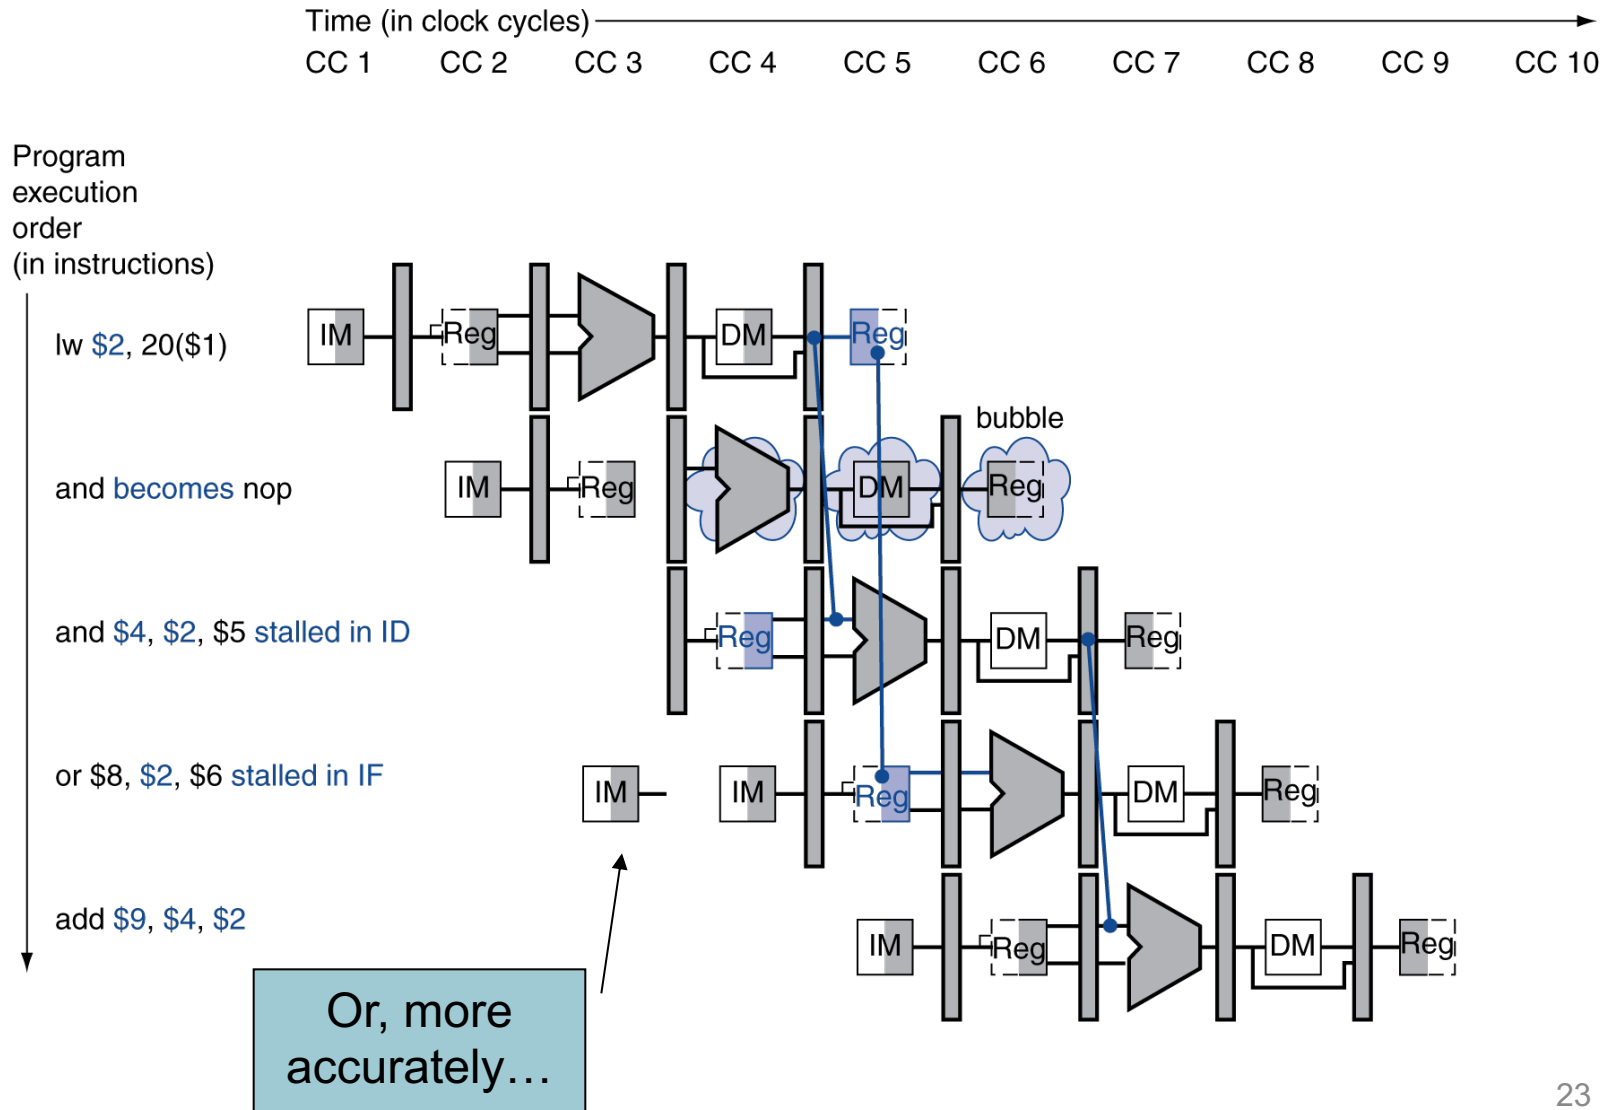
How to Stall the Pipeline

- **Force control values in ID/EX register to 0**
 - EX, MEM and WB do a **no-operation** (nop)
- **Prevent update of PC and IF/ID register**
 - Using instruction is decoded again
 - Following instruction is fetched again
 - 1-cycle stall allows MEM to read data for 1w
 - Can subsequently forward to EX stage

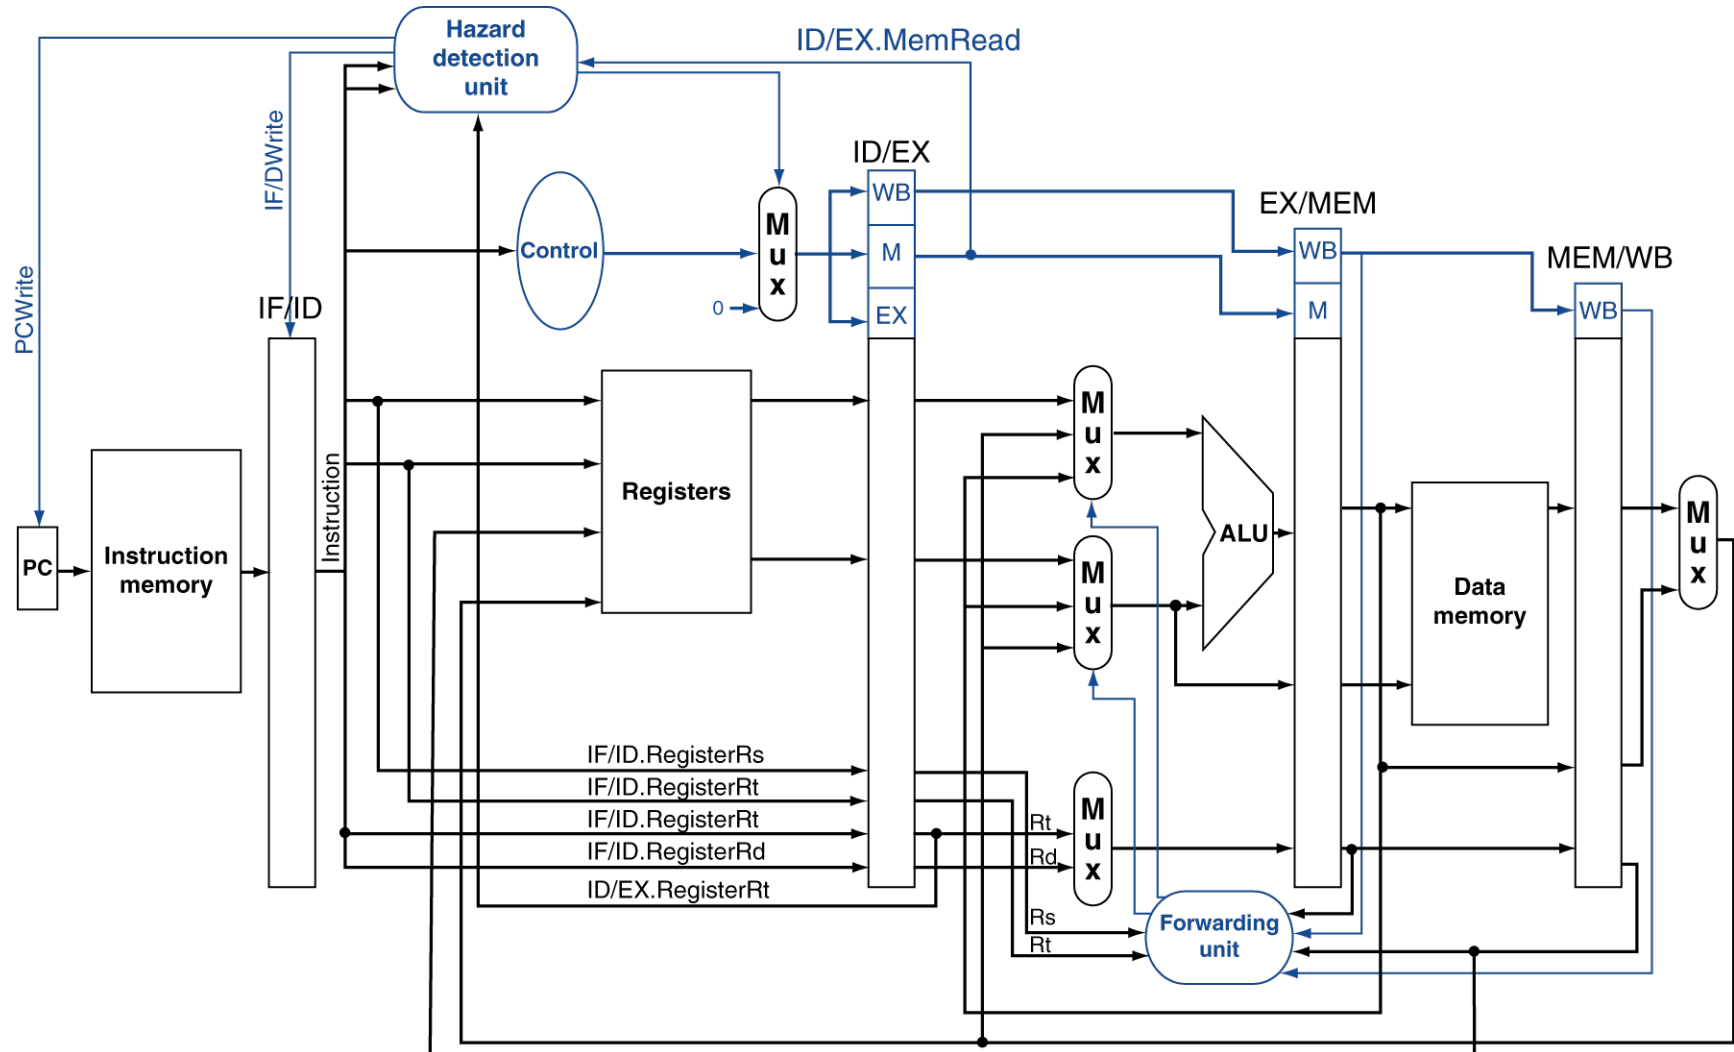
Load-Use Data Hazard



Load-Use Data Hazard



Datapath with Hazard Detection

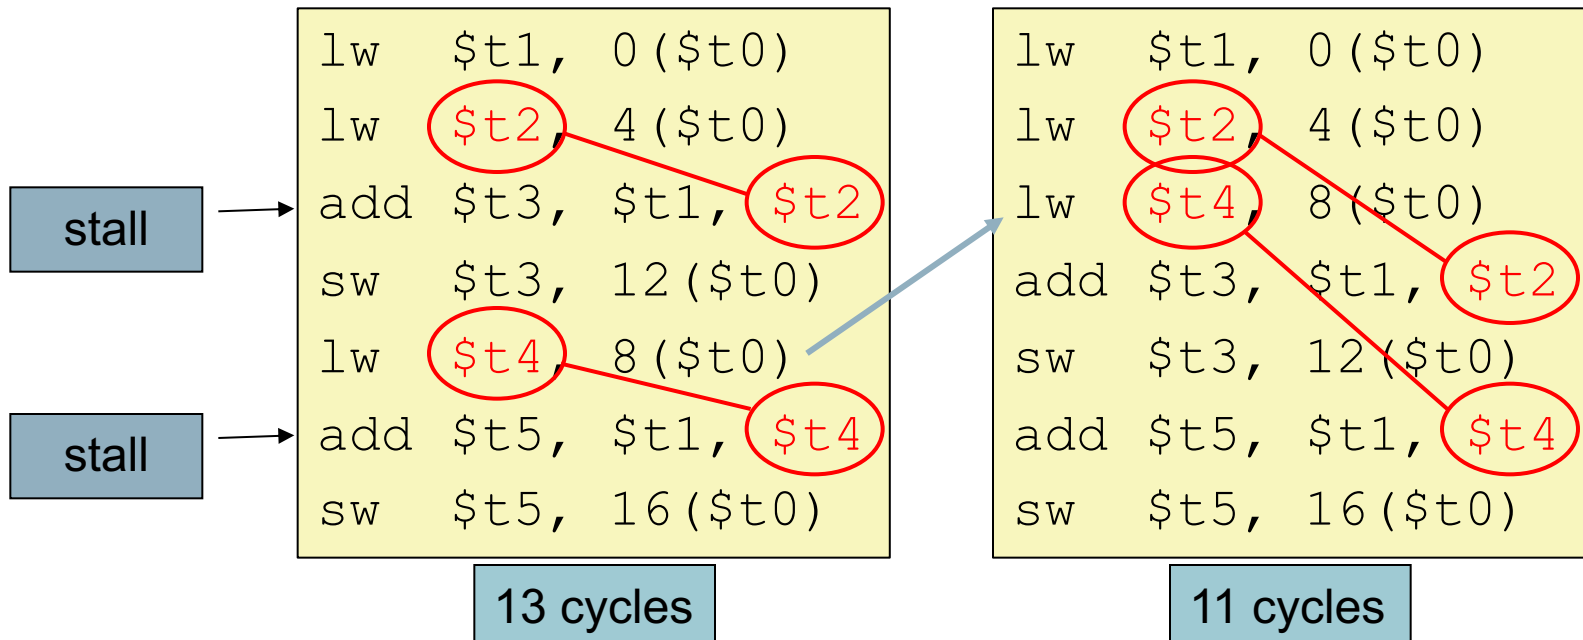


Stalls and Performance

- **Stalls reduce performance but are required to get correct results**
- **Compiler can arrange code to avoid hazards and stalls**
 - Although it generally relies upon the hardware to resolve hazards and thereby ensure correct execution
 - Requires knowledge of the pipeline structure

Code Reordering to Avoid Stalls

- Reorder code to avoid use of load result in the next instruction
- Consider the code for $A = B + E$; $C = B + F$;



Data hazards: Summary

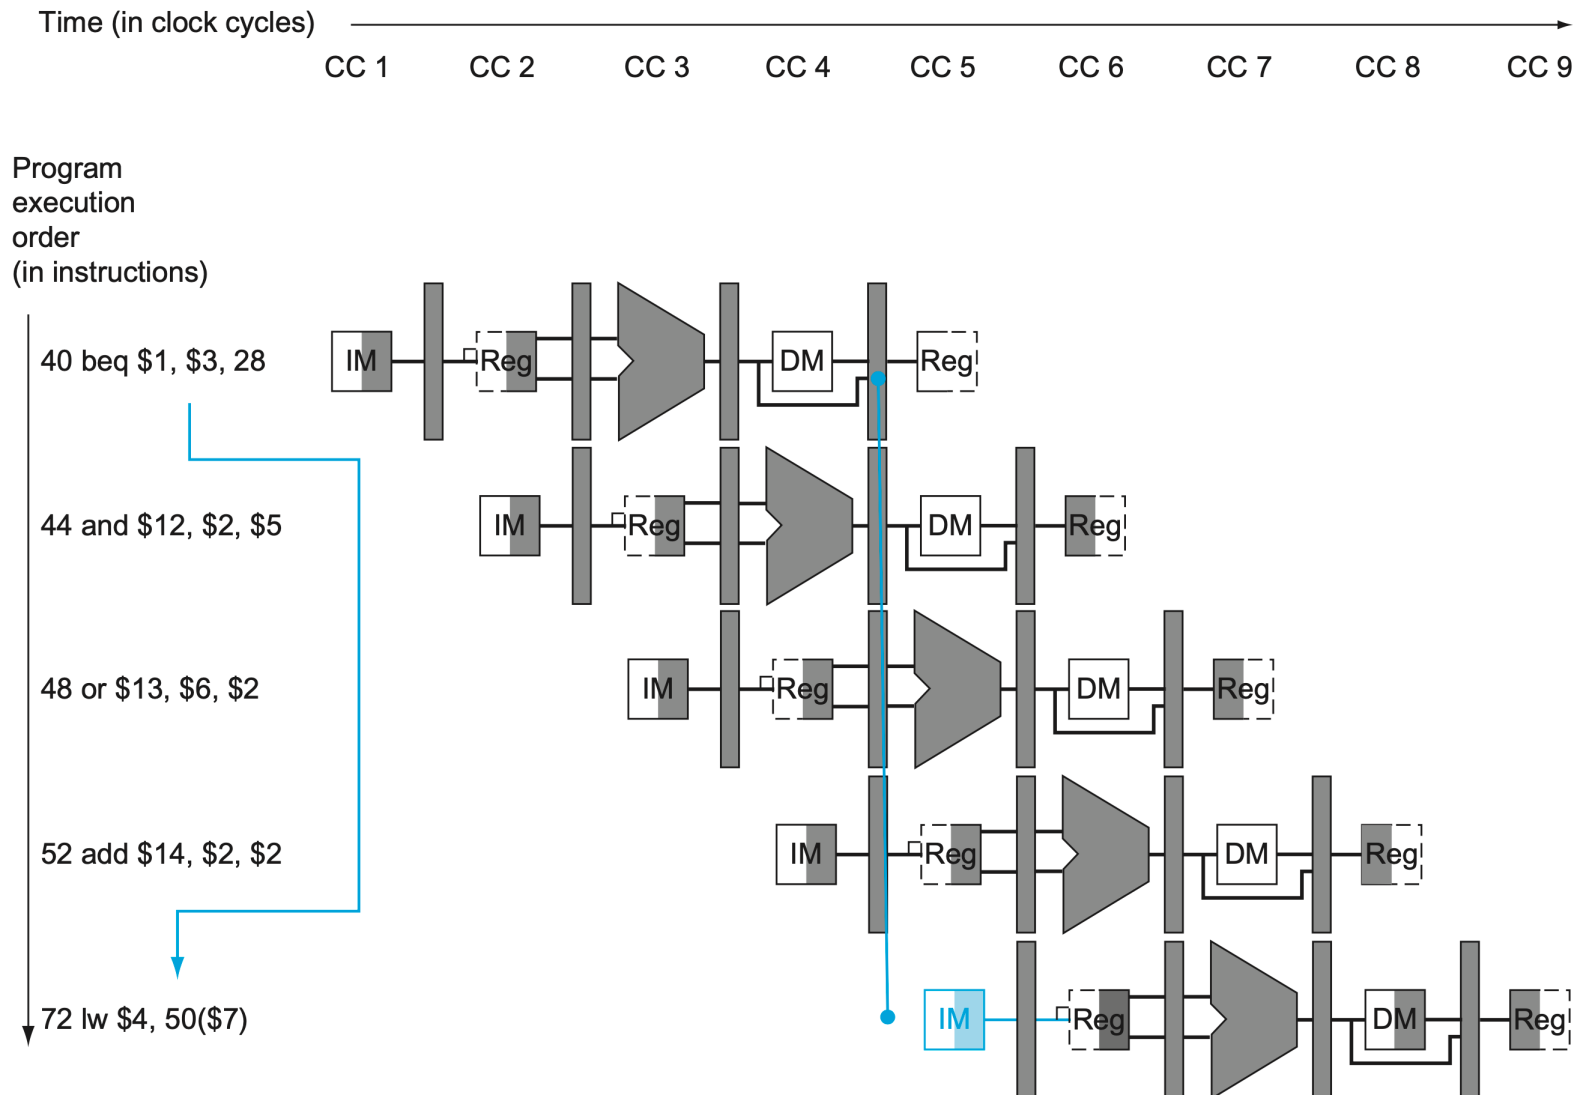
Situation	Example	Action
No dependence	<pre>lw \$t1, 45(\$t2) add \$t5, \$t6, \$t7 sub \$t8, \$t6, \$t7 or \$t9, \$t6, \$t7</pre>	No hazard possible because no dependence exists on \$t1 in the immediately following three instructions
Dependence requiring stall	<pre>lw \$t1, 45(\$t2) add \$t5, \$t1, \$t7 sub \$t8, \$t6, \$t7 or \$t9, \$t6, \$t7</pre>	Comparators detect the use of \$t1 in the ADD and stall the ADD (and SUB and OR) before the ADD begins EX
Dependence overcome by forwarding	<pre>lw \$t1, 45(\$t2) add \$t5, \$t6, \$t7 sub \$t8, \$t1, \$t7 or \$t9, \$t6, \$t7</pre>	Comparators detect the use of \$t1 in SUB and forward the result of LOAD to the ALU in time for SUB to begin with EX
Dependence with accesses in order (read after write)	<pre>lw \$t1, 45(\$t2) add \$t5, \$t6, \$t7 sub \$t8, \$t6, \$t7 or \$t9, \$t1, \$t7</pre>	No action is required because the read of \$t1 by OR occurs in the second half of the ID phase, while the write of the loaded data occurred in the first half

Control Hazards

- **An instruction must be fetched at every clock cycle to sustain the pipeline**
 - Yet, in our design, the decision about whether to branch doesn't occur until the MEM pipeline stage
- **Branch determines flow of control**
 - Fetching next instruction depends on branch outcome
 - Pipeline **can't always fetch correct instruction**
 - Still working on ID stage of branch

```
36:  sub    $10, $4, $8
40:  beq   $1,  $3, 7      # PC-relative addressing
44:  and    $12, $2, $5
48:  or     $13, $2, $6
52:  add    $14, $4, $2
56:  slt    $15, $6, $7
    ...
72:  lw     $4, 50($7)
```

Control Hazards



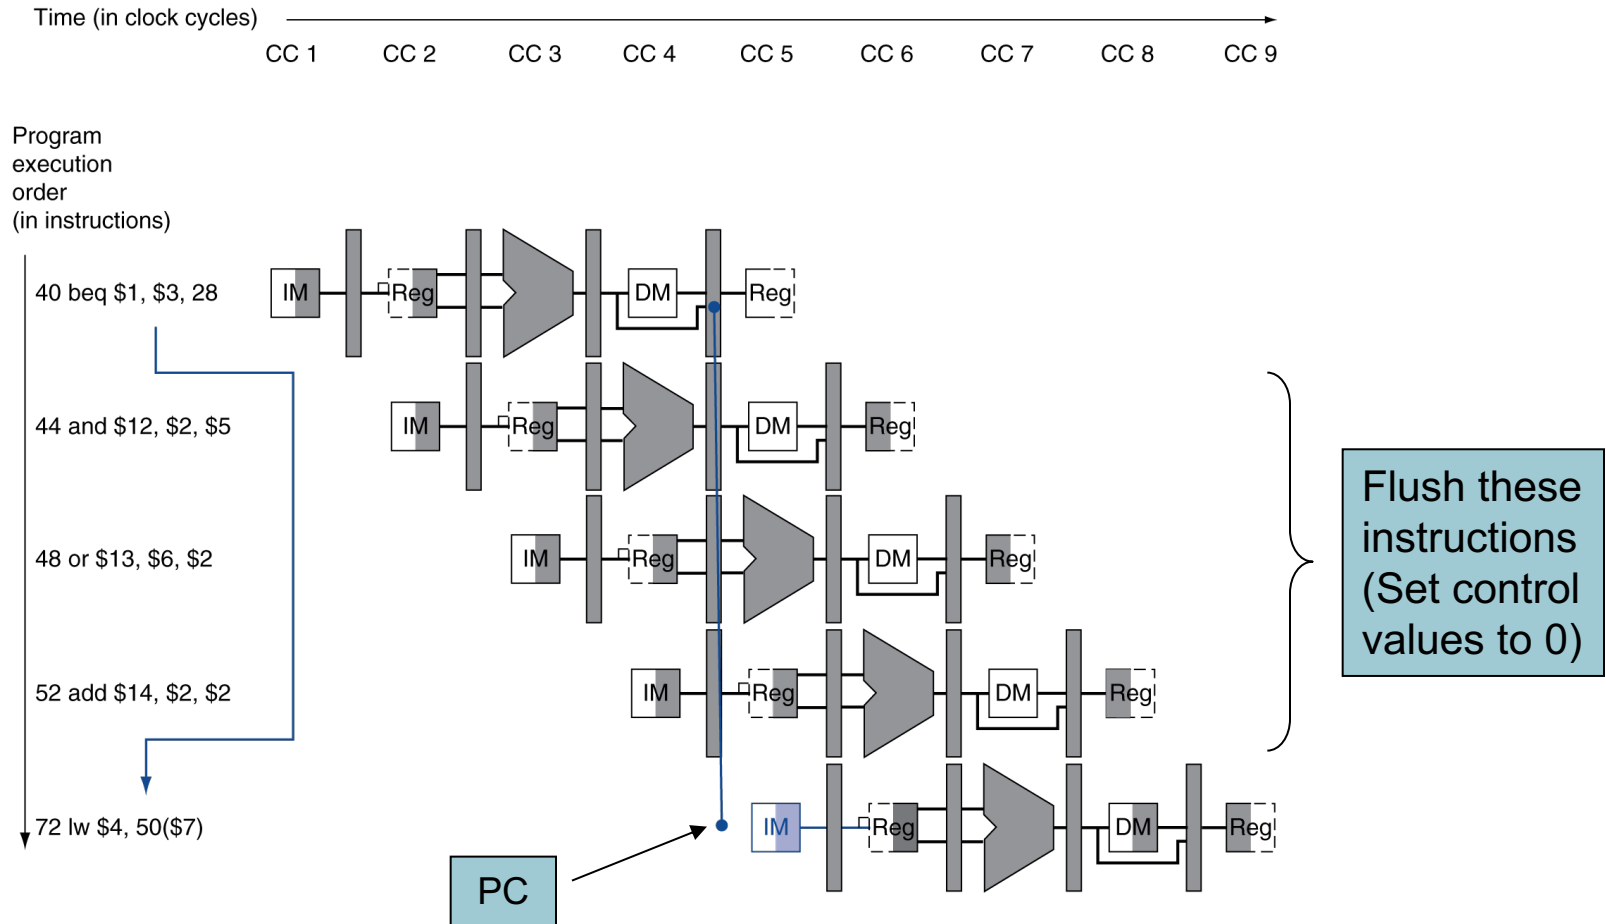
Stall on Branch

- **Wait until branch outcome is determined before fetching next instruction**
 - Makes all branches cost 4 cycles
- **Seems wasteful, particularly when the branch isn't taken**
 - Could have spent clock cycles fetching, decoding and executing next instructions
- **The cost of this option is too high for most computers to use and motivates another solutions**

Assume Branch Not Taken

- **Always assume branch will NOT be taken**
 - On average, branches are taken half the time
- **Cuts overall time for branch processing in half**
 - If prediction is correct, continue execution down the sequential instruction stream
 - If prediction is incorrect, just discard the instructions
- **Discarding instructions means we must be able to flush instructions in the IF, ID, and EX stages**
 - Just change control to 0

Assume Branch Not Taken

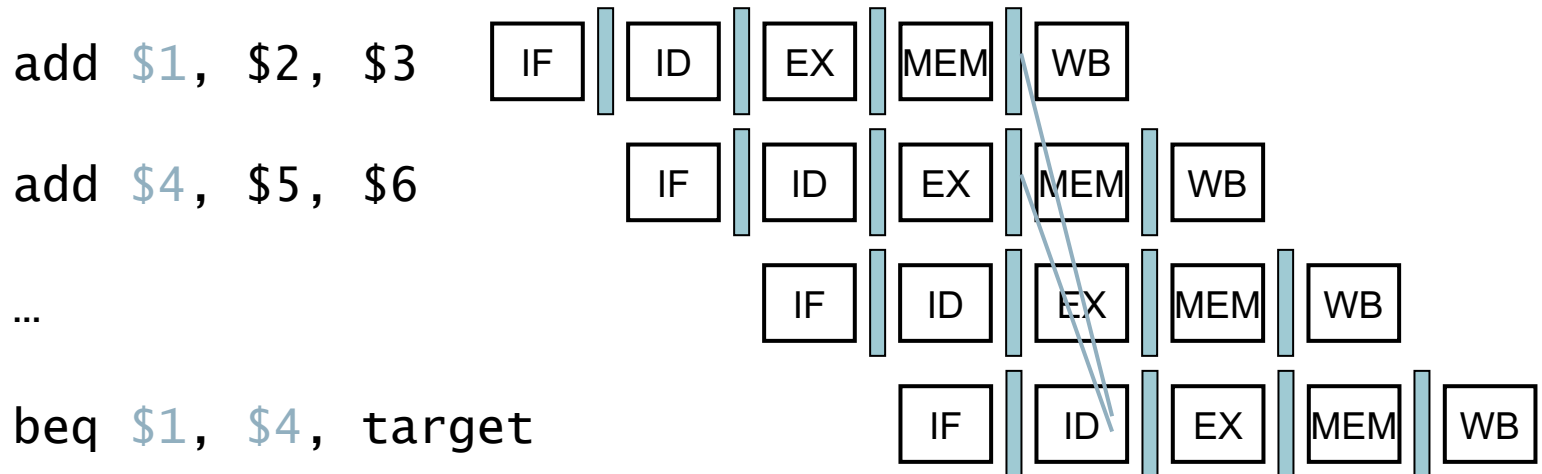


Reducing Branch Delay

- **Move hardware to determine outcome to ID stage**
 - Reduces the **penalty of a branch to only one instruction** if the branch is taken, namely, the one currently being fetched
- **Branch address calculation (easy part)**
 - We already have the PC value and the immediate field in the IF/ID pipeline register
 - Just move the branch adder from the EX stage to the ID stage
- **Moving the branch test implies additional forwarding and hazard detection hardware**
 - Since a branch dependent on a result still in the pipeline must still work properly with this optimization

Data Hazards for Branches

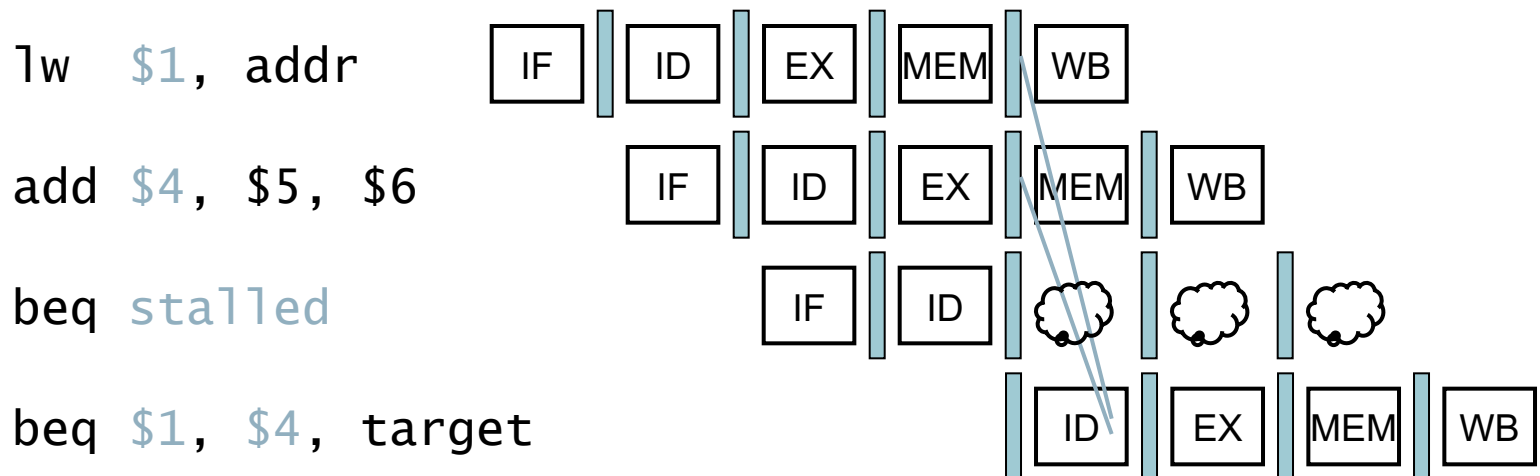
- If a comparison register is a destination of 2nd or 3rd preceding ALU instruction



- Can resolve using forwarding

Data Hazards for Branches


- If a comparison register is a destination of preceding ALU instruction or 2nd preceding load instruction



- Need 1 stall cycle

Data Hazards for Branches

- If a comparison register is a destination of immediately preceding load instruction

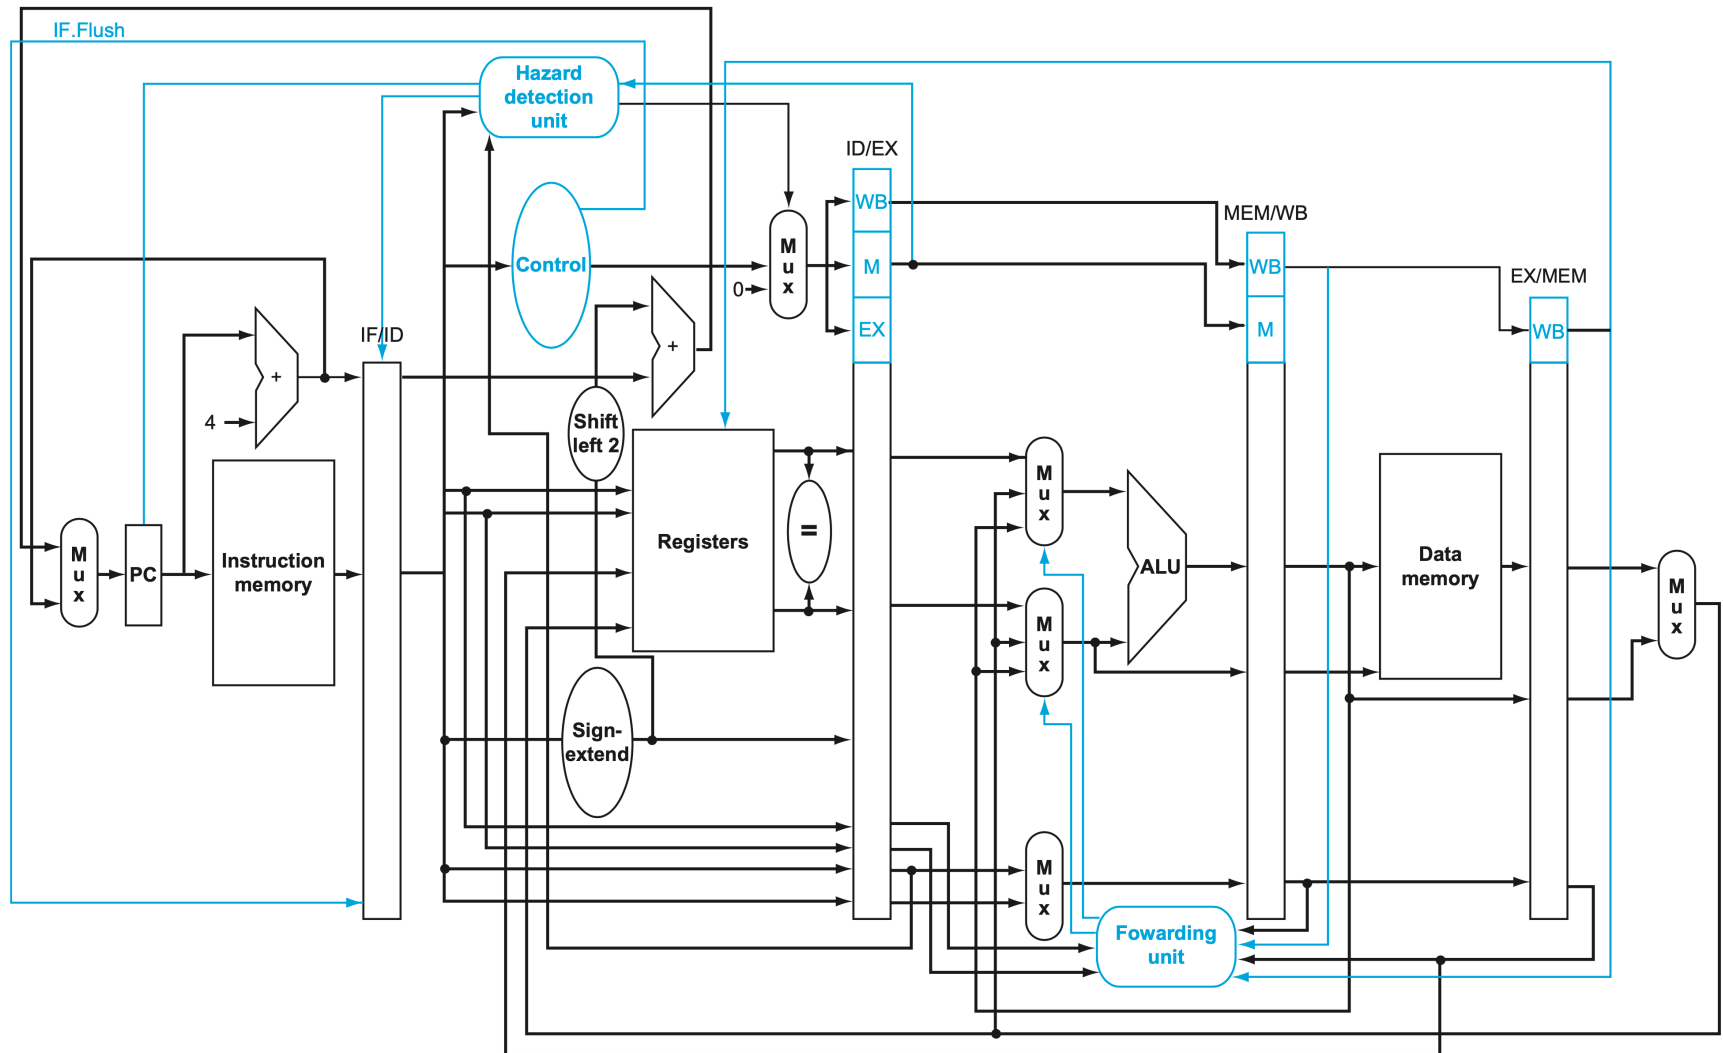


- Need 2 stall cycles

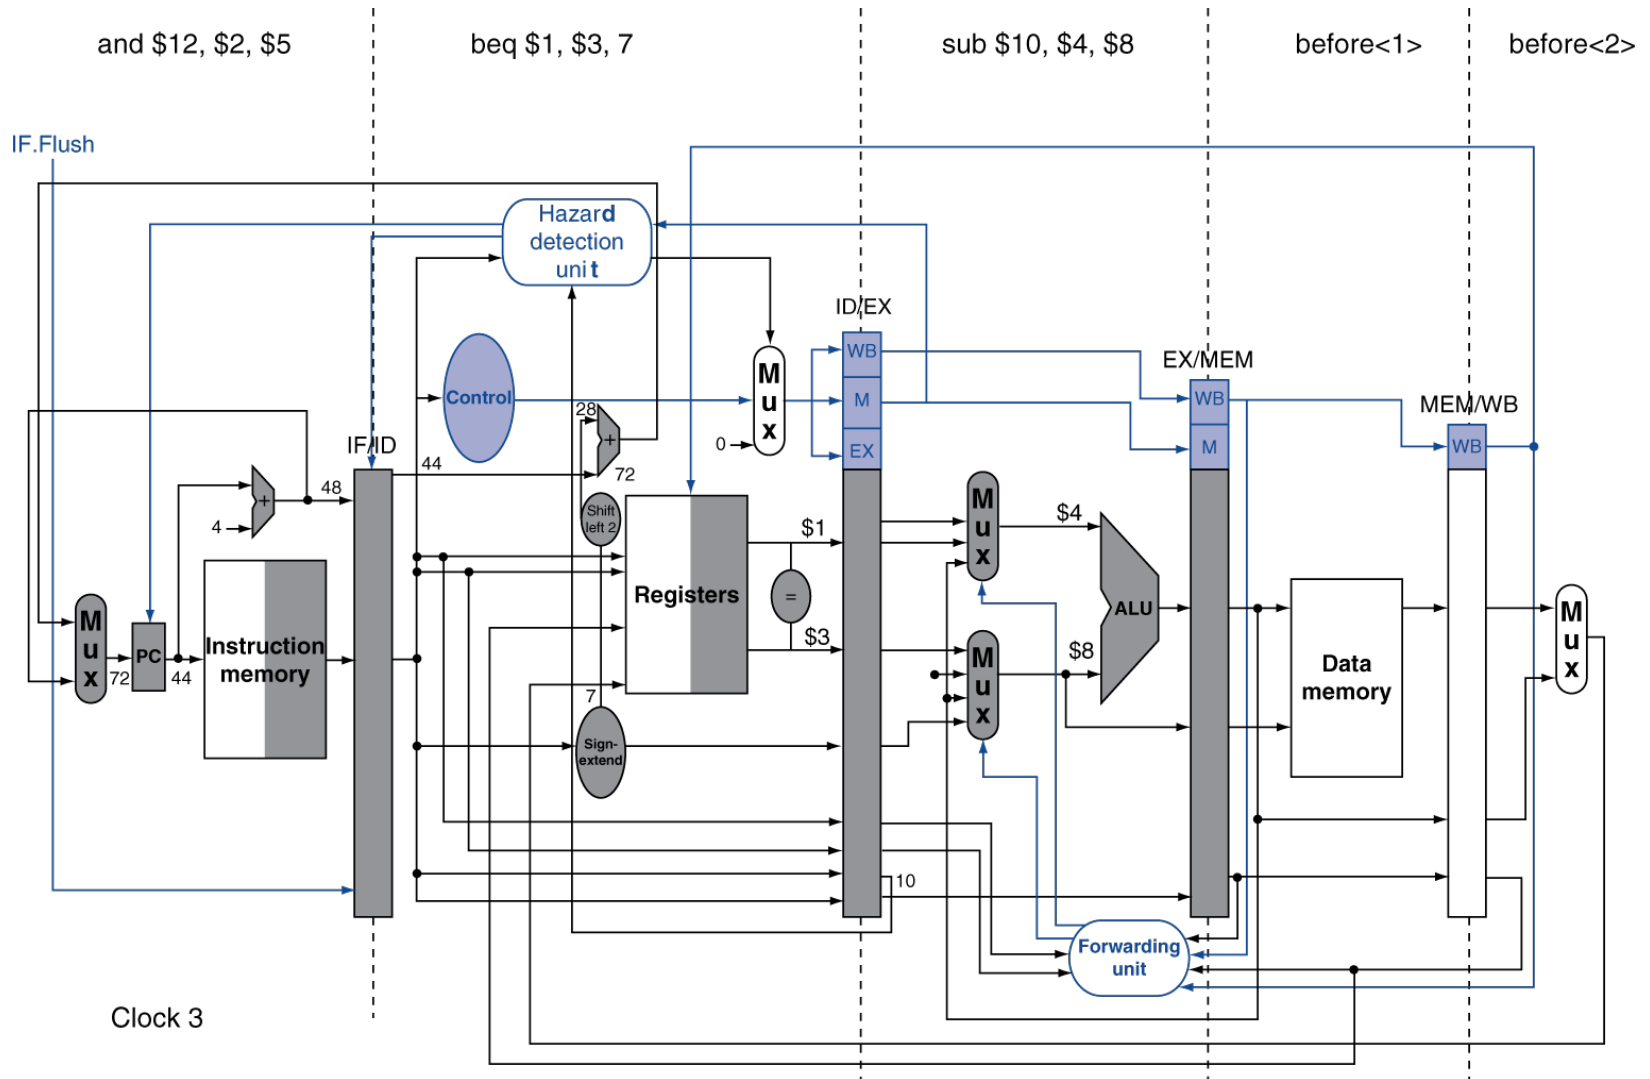
Reducing Branch Delay

- **Despite these difficulties, moving the branch execution to the ID stage **is an improvement****
 - Reduces the penalty of a branch to only one instruction if the branch is taken, namely, the one currently being fetched
- **To flush instructions in the IF stage, we add a control line, called IF.Flush**
 - That zeros the instruction field of the IF/ID pipeline register
 - Clearing the register transforms the fetched instruction into a `nop`

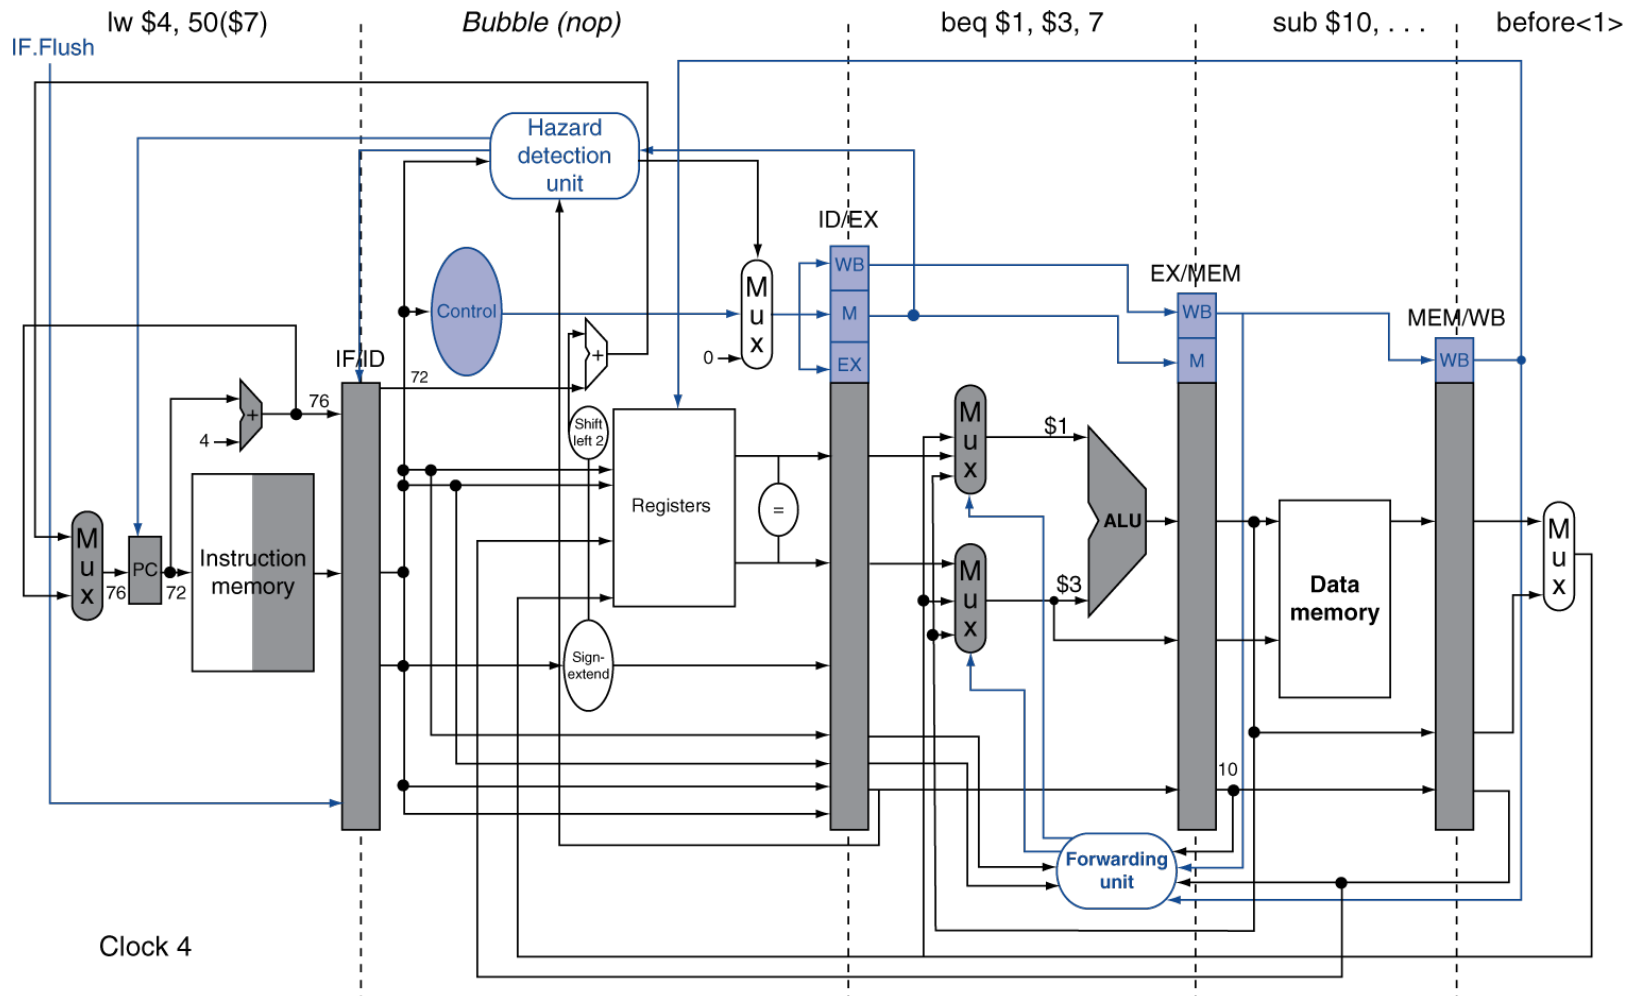
Reducing Branch Delay



Example: Branch Taken

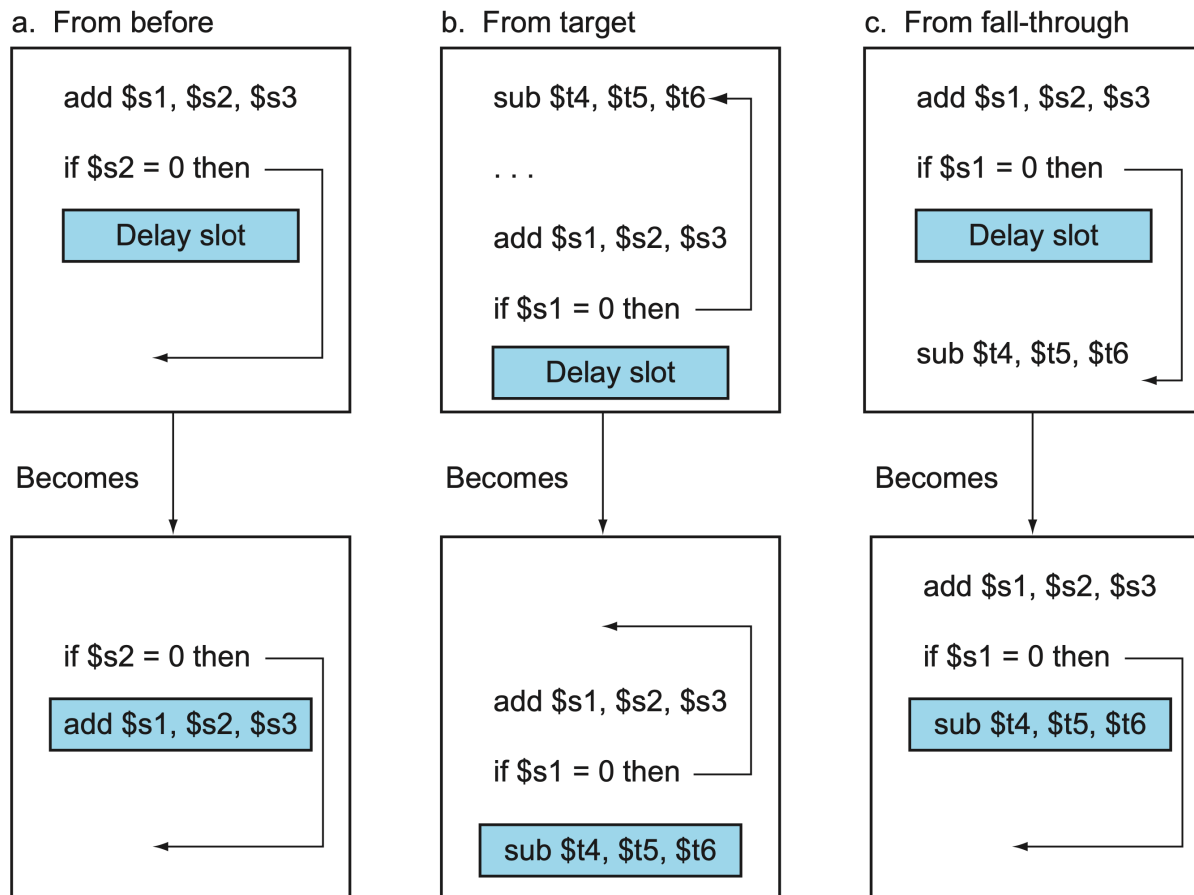


Example: Branch Taken



Branch Delay Slot

- Compiler tries to place an instruction that always executes after the branch in the **branch delay slot**
 - b) and c) are used when a) is not possible



Branch Delay Slot

- **Limitations on delayed branch scheduling arise from**
 1. The restrictions on the instructions that are scheduled into the delay slots
 2. The ability to predict at compile time whether a branch is likely to be taken or not

- **Is a simple and effective solution for a five-stage pipeline issuing one instruction each clock cycle**
 - As processors go to longer pipelines, the branch delay becomes longer, and a single delay slot is insufficient
 - Hence, delayed branching has lost popularity compared to more expensive but more flexible dynamic approaches

Concluding Remarks

- **Pipelining improves instruction throughput using parallelism**
 - More instructions completed per second
 - Latency for each instruction not reduced
- **There are situations when the next instruction cannot execute in the following clock cycle**
 - These events are called hazards
- **Structural hazards**
 - Occur when the hardware cannot support the combination of instructions that we want to execute in the same clock cycle
 - Avoided by replicating functional units

Concluding Remarks

■ Data hazards

- Occur when the pipeline must be stalled because one step must wait for another to complete
- Partially solved through forwarding
 - Demands adding extra hardware to retrieve the missing item early from the internal resources

■ Control hazards

- Arising from the need to decide based on the results of one instruction while others are executing
- Partially solved through reducing the branch delay and prediction
 - Demands moving hardware to determine outcome of branch to ID stage and adding hardware to support prediction